Computer Organization and Architecture

Module 5 (Part 2)
Design of Memory Subsystems

Prof. Indranil Sengupta

Dr. Sarani Bhattacharya

Department of Computer Science and Engineering

IIT Kharagpur
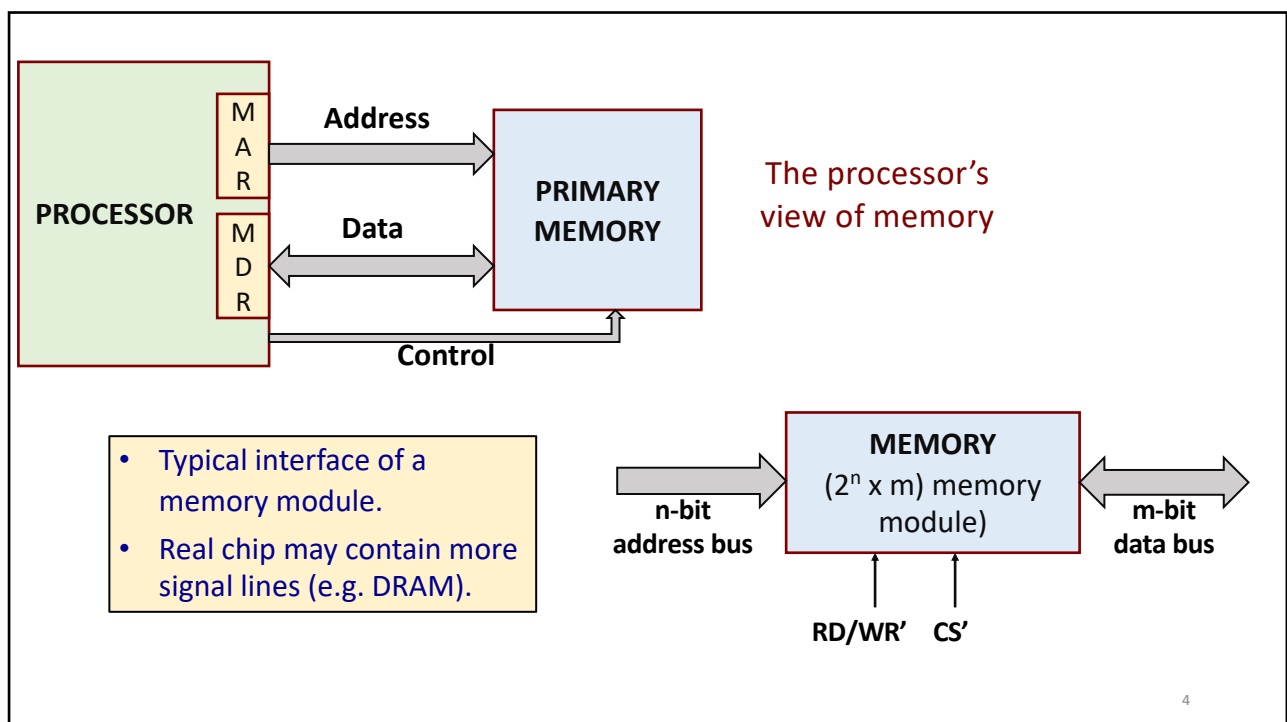
1

Memory Interfacing and Addressing

2

Memory Interfacing

- Basic problem:
  - Interfacing one of more memory modules to the processor.
  - We assume a single level memory at present (i.e. no cache memory).
- Questions to be answered:
  - How the processor address and data lines are connected to memory modules?
  - How are the addresses decoded?
  - How are the memory addresses distributed among the memory modules?
  - How to speed up data transfer rate between processor and memory?

3

3

The processor's view of memory

- Typical interface of a memory module.
- Real chip may contain more signal lines (e.g. DRAM).

MEMORY
(2^n x m) memory module)

n-bit address bus

m-bit data bus

RD/WR'   CS'

4

4

A Note About the Memory Interface Signals

- The data signals of a memory module (RAM) are typically bidirectional.
  - Some memory chips may have separate data in and data out lines.

- For memory *READ* operation:
  - Address of memory location is applied to *address lines*.
  - *RD/WR'* control signal is set to 1, and *CS'* is set to 0.
  - Data is read out through the *data lines* after memory access time delay.

- For memory *WRITE* operation:
  - Address of memory location is applied to *address lines*, and the data to be written to *data lines*.
  - *RD/WR'* control signal is set to 0, and *CS'* is set to 0.

5

5

- Why is *CS'* signal required?
  - To handle multiple memory modules interfacing problem.
  - We typically select only one out of several memory modules at a time.

- What happens when *CS' = 1*?
  - When a memory module is *not selected*, the data lines are set to the *high impedance state* (i.e. electrically disconnected).
  - An example scenario is shown.

SELECTED

Address → **MEMORY** $(2^n \times m)$

RD/WR'   CS' = 0

Address → **MEMORY** $(2^n \times m)$

RD/WR'   CS' = 1

**m-bit data bus**

6

6

3

An Example Memory Interfacing Problem

- Consider a MIPS32 like processor with a 32-bit address.
  - Maximum memory that can be connected is 2^{32} = 4 Gbytes.
  - Assume that the processor data lines are 8 bits.

- Assume that memory chips (RAM) are available with *size 1 Gbyte*.
  - 30 address lines and 8 data lines.
  - Low-order 30 address lines (A_{29}-A_0) are connected to the memory modules.

- We want to interface *4 such chips* to the processor.
  - Total memory of 4 Gbytes.

7

7

8

8

4
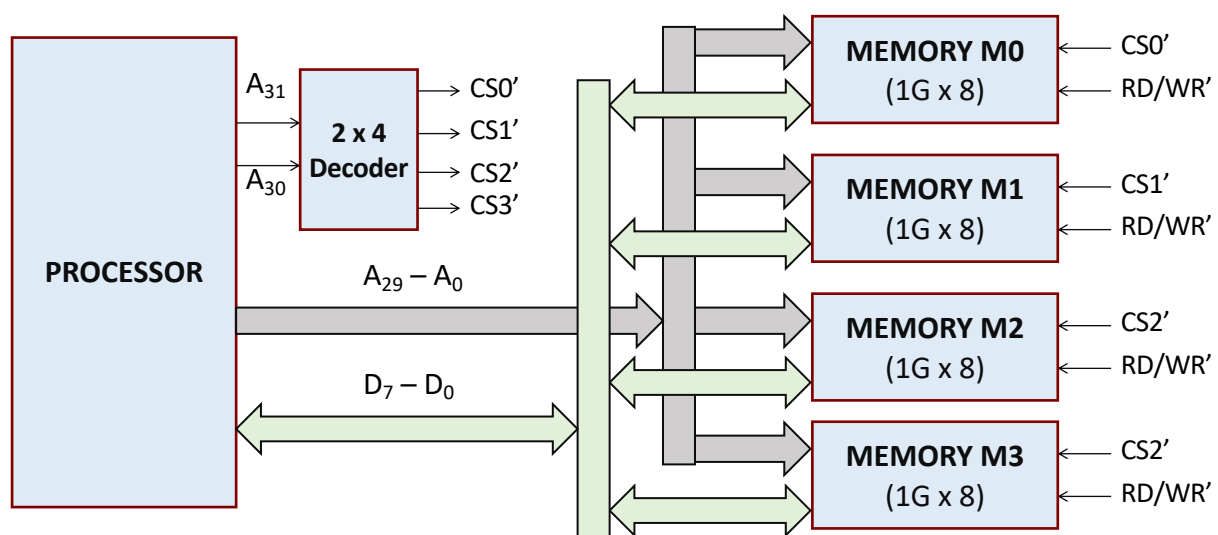
- High order address lines (A_{31} and A_{30}) select one of the memory modules.
- **When is M0 selected?**
  - Address is: 0 0 x x x x x x x x x x x x x x x x x x x x x x x x x x x x x x
  - Range of addresses is: 0x00000000 to 0x3FFFFFFF
- **When is M1 selected?**
  - Address is: 0 1 x x x x x x x x x x x x x x x x x x x x x x x x x x x x x x
  - Range of addresses is: 0x40000000 to 0x7FFFFFFF
- **When is M2 selected?**
  - Address is: 1 0 x x x x x x x x x x x x x x x x x x x x x x x x x x x x x x
  - Range of addresses is: 0x80000000 to 0xBFFFFFFF
- **When is M3 selected?**
  - Address is: 1 1 x x x x x x x x x x x x x x x x x x x x x x x x x x x x x x
  - Range of addresses is: 0xC0000000 to 0xFFFFFFFF

9

9

- **An observation:**
  - Consecutive block of bytes are mapped to the same memory module.
  - For MIPS32, we have to access 32 bits (4 bytes) of data in parallel, which requires four sequential memory accesses here.
  - We shall look at an alternate memory organization later that would make this possible.
   - Called *memory interleaving*.

10

10

Exercise 1

• 6116 is a 2K x 8 RAM chip. Build a 16K x 8 memory using 6116 chips. Show all the connections, and state how the addresses are distributed across memory modules.

6116 Pinout

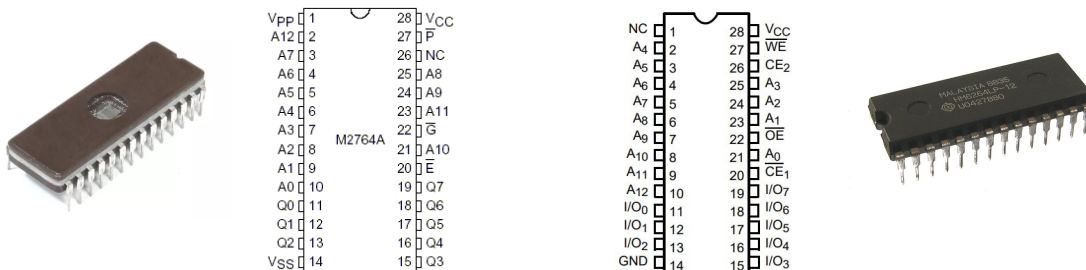|			
A_7	1	24	VCC
A_6	2	23	A_8
A_5	3	22	A_9
A_4	4	21	$\overline{WE}$
A_3	5	20	$\overline{OE}$
A_2	6	19	A_{10}
A_1	7	18	$\overline{CS}$
A_0	8	17	I/O_7
I/O_0	9	16	I/O_6
I/O_1	10	15	I/O_5
I/O_2	11	14	I/O_4
GND	12	13	I/O_3

(6116)

11

Exercise 2

• 6264 is an 8K x 8 RAM chip, and 2764 is an 8K x 8 EPROM chip. Build a 32K x 8 memory system with 8 KB EPROM and 24 KB RAM. Show all the connections, and state how the addresses are distributed across memory modules.

6264 Pinout

M2764A pinout:

V_{PP}	1	28	V_{CC}
A_{12}	2	27	$\overline{P}$
A_7	3	26	NC
A_6	4	25	A_8
A_5	5	24	A_9
A_4	6	23	A_{11}
A_3	7	22	$\overline{G}$
A_2	8	21	A_{10}
A_1	9	20	$\overline{E}$
A_0	10	19	Q7
Q0	11	18	Q6
Q1	12	17	Q5
Q2	13	16	Q4
V_{SS}	14	15	Q3

6264 pinout:

|			
NC	1	28	V_{CC}
A_4	2	27	$\overline{WE}$
A_5	3	26	CE_2
A_6	4	25	A_3
A_7	5	24	A_2
A_8	6	23	A_1
A_9	7	22	$\overline{OE}$
A_{10}	8	21	A_0
A_{11}	9	20	$\overline{CE_1}$
A_{12}	10	19	I/O_7
I/O_0	11	18	I/O_6
I/O_1	12	17	I/O_5
I/O_2	13	16	I/O_4
GND	14	15	I/O_3

12

6

Exercise 3

- 2114 is a 1K x 4 RAM chip. Build a 4K x 8 memory using 2114 chips. Show all the connections, and state how the addresses are distributed across memory modules.

2114N Pinout

13

Memory Interleaving

14

Improved Memory Interface for MIPS32

- We make small changes in the organization so that 32-bits of data can be fetched in a single memory access cycle.
  - Exploit the concept of *memory interleaving*.
  - Consecutive bytes are mapped to different memory modules.
- The main changes:
  - High order 30 address lines (A_{31}-A_2) are connected to memory modules.
  - Low order two address lines (A_1 and A_0) are used to select one of the modules.

15

15

- How are the addresses mapped to memory modules?
  - *Module M0*: 0, 4, 8, 12, 16, 20, 24, …
  - *Module M1*: 1, 5, 9, 13, 17, 21, 25, …
  - *Module M2*: 2, 6, 10, 14, 18, 22, 26, …
  - *Module M3*: 3, 7, 11, 15, 19, 23, 27, …
- Memory addresses are *interleaved* across memory modules.
- What we can gain from this mapping?
  - Consecutive addresses are mapped to consecutive modules.
  - Possible to access four consecutive words in the same cycle, if all four modules are enabled simultaneously.
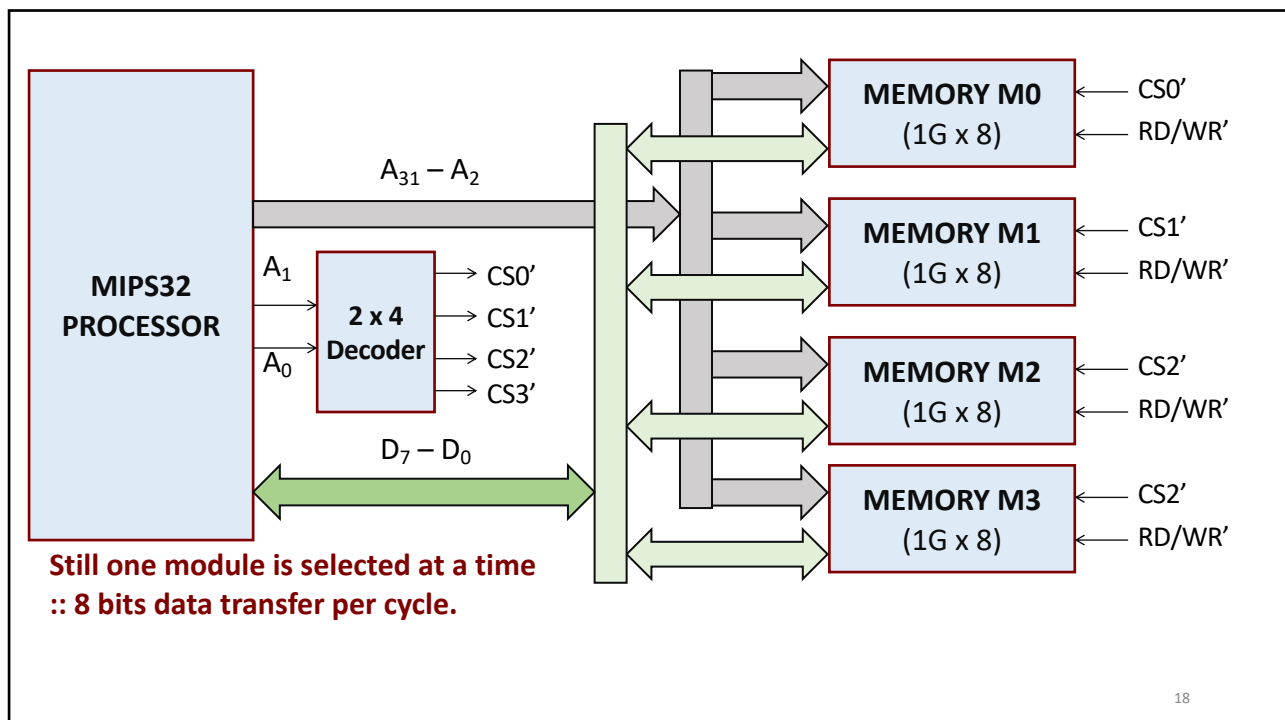
16

16

- Motivation for word alignment in MIPS32 data words.
  - 32-bit words start from a memory address that is divisible by 4.
    - Corresponding byte addresses are (0, 1, 2, 3), (4, 5, 6, 7), (8, 9, 10, 11), (12, 13, 14, 15), etc. → *last two bits of the addresses are 00, 01, 10 and 11*.
    - Possible to transfer all the four bytes of a word in a *single memory cycle*.
  - What happens if a word is not aligned?
    - Say: (1, 2, 3, 4) or (2, 3, 4, 5) or (3, 4, 5, 6).   **2 memory cycles required**
    - Two of the bytes will be mapped to the same memory module.
    - Hence the word cannot be transferred in a single memory cycle.

17

17

MIPS32 PROCESSOR

$A_{31} - A_2$

A_1

2 x 4 Decoder

A_0

CS0'
CS1'
CS2'
CS3'

$D_7 - D_0$

MEMORY M0
(1G x 8)

CS0'
RD/WR'

MEMORY M1
(1G x 8)

CS1'
RD/WR'

MEMORY M2
(1G x 8)

CS2'
RD/WR'

MEMORY M3
(1G x 8)

CS2'
RD/WR'

Still one module is selected at a time :: 8 bits data transfer per cycle.

18

18

9

Enable all the four modules together.
32-bit parallel data transfer.
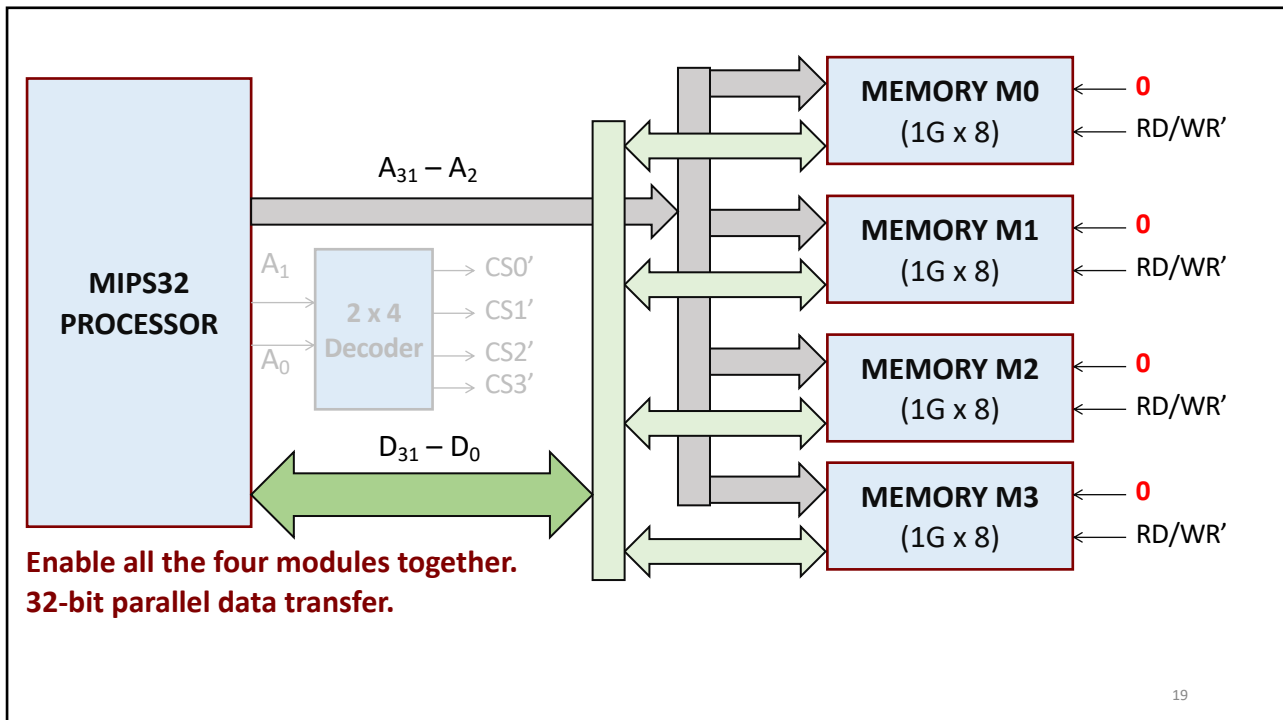
19

Memory Latency and Bandwidth

- **Memory Latency:**
  - The delay from the issue of a memory read request to the first byte of data becoming available.

- **Memory Bandwidth:**
  - The maximum number of bytes that can be transferred between the processor and the memory system per unit time.

20

- **Example 1:**

  Consider a memory system that takes 20 ns to service the access of a single 32-bit word.

  Latency L = 20 ns per 32-bit word.

  Bandwidth BW = 32 / (20 x 10^{-9}) = 200 Mbytes per second.

- **Example 2:**

  The memory system is modified to accept a new (still 20ns) request for a 32-bit word every 5 ns by overlapping requests.

  Latency L = 20 ns per 32-bit word (*no change*).

  Bandwidth BW = 32 / (5 x 10^{-9}) = 800 Mbytes per second.

21

21

Memory Hierarchy Design

22

Introduction

- Programmers want unlimited amount of memory with very low latency.

- Fast memory technology is more expensive per bit than slower memory.
  - SRAM is more expensive than DRAM, DRAM is more expensive than disk.

- Possible solution?
  - Organize the memory system in several levels, called *memory hierarchy*.
  - Exploit *temporal locality* and *spatial locality* of computer programs.
  - Try to keep the commonly accessed segments of program / data in the faster memories.
  - Results in faster access times on the average.

23

23

Quick Review of Memory Technology

- **Static RAM:**
  - Very fast but expensive memory technology (requires 6 transistors / bit).
  - Packing density is limited.
- **Dynamic RAM:**
  - Significantly slower than SRAM, but much less expensive (1 transistor / bit).
  - Requires periodic refreshing.
- **Flash memory / Solid-state drive:**
  - Non-volatile memory technology that uses floating-gate MOS transistors.
  - Slower than DRAM, but higher packing density, and lower cost per bit.
- **Magnetic disk:**
  - Provides large amount of storage, with very low cost per bit.
  - Much slower than DRAM, and also flash memory.
  - Requires mechanical moving parts, and uses magnetic recording technology.

24

24

Memory Hierarchy

- The memory system is organized in *several levels*, using progressively faster technologies as we move towards the processor.
    - The entire addressable memory space is available in the largest (but slowest) memory (typically, magnetic disk or flash storage).
    - We incrementally add smaller (but faster) memories, each containing a *subset* of the data stored in the memory below it.
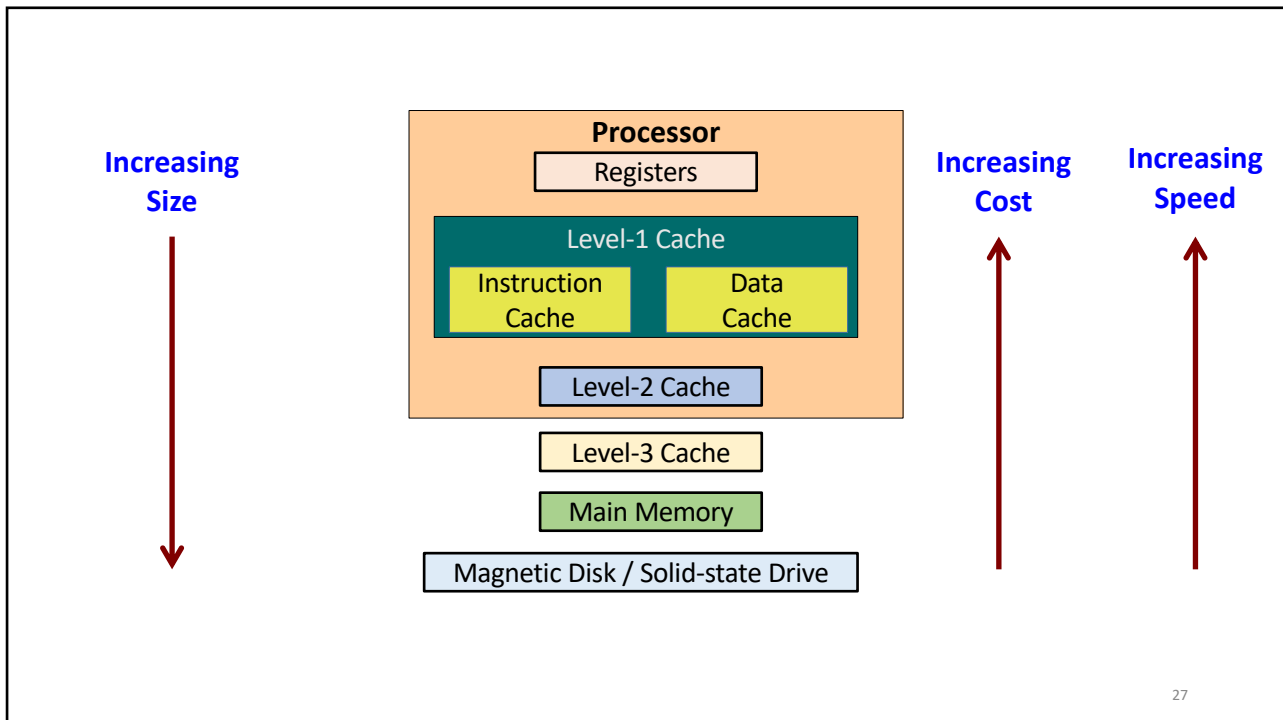   - We proceed in steps towards the processor.

25

25

- Typical hierarchy (*starting with the one closest to the processor*):
    1. Processor registers
    2. Level-1 cache (typically divided into separate instruction and data cache)
    3. Level-2 cache
    4. Level-3 cache
    5. Main memory
    6. Secondary memory (magnetic disk / solid-state drive)
- As we move away from the processor:
    a) Size increases
   b) Cost per bit decreases
    c) Speed decreases

26

26

Increasing Size

Processor
Registers

Level-1 Cache

| Instruction Cache | Data Cache |

Level-2 Cache

Level-3 Cache

Main Memory

Magnetic Disk / Solid-state Drive

Increasing Cost

Increasing Speed

27

A Comparison

Level	Typical Access Time	Typical Capacity	Other Features
Register	300-500 ps	500-1000 B	On-chip
Level-1 cache	1-2 ns	16-64 KB	On-chip
Level-2 cache	5-20 ns	256 KB – 2 MB	On-chip
Level-3 cache	20-50 ns	1-32 MB	On or off chip
Main memory	50-100 ns	1-16 GB	Off-chip
Magnetic disk	5-50 ms	100 GB – 16 TB	

28

Processor-Memory Performance Gap

- Processor is much faster than main memory.
  - Processor has to spend much of the time waiting while instructions and data are being fetched from main memory.
  - Memory speed cannot be increased beyond a certain point.

Impact of Processor / Memory Performance Gap

Year	CPU Clock	Clock Cycle	Memory Access	Minimum CPU Stall Cycles
1986	8 MHz	125 ns	190 ns	190 / 125 − 1 = 0.5
1989	33 MHz	30 ns	165 ns	165 / 30 − 1 = 4.5
1992	60 MHz	16.6 ns	120 ns	120 / 16.6 − 1 = 6.2
1996	200 MHz	5 ns	110 ns	110 / 5 − 1 = 21.0
1998	300 MHz	3.33 ns	100 ns	100 / 3.33 − 1 = 29.0
2000	1 GHz	1 ns	90 ns	90 / 1 − 1 = 89.0
2002	2 GHz	0.5 ns	80 ns	80 / 0.5 − 1 = 159.0
2004	3 GHz	0.33 ns	60 ns	60 / 0.33 − 1 = 179.0

Ideal memory access time = 1 CPU cycle

Real memory access time >> 1 CPU cycles

- **Memory Latency Reduction Techniques:**
  - Faster DRAM cells (depends on VLSI technology)
  - Wider memory bus width (fewer memory accesses needed)
  - Multiple memory banks
  - Integration of memory controller with processor
  - New emerging RAM technologies
- **Memory Latency Hiding Techniques**
  - Memory hierarchy (using SRAM-based cache memories)
  - Pre-fetching instructions and/or data from memory before they are actually needed (used to hide long memory access latency)

31

31

Locality of Reference

- Programs tend to reuse data and instructions they have used recently.
  - **Rule of thumb**: 90% of the total execution time of a program is spent in only 10% of the code (also called *90/10 rule*).
  - **Reason**: nested loops in a program, few procedures calling each other repeatedly, arrays of data items being accessed sequentially, etc.
- Basic idea to exploit this rule:
  - Based on a program's recent past, we can predict with a reasonable accuracy what instructions and data will be accessed in the near future.

32

32

- The *90/10 rule* has two dimensions:

  a) **Temporal Locality** (locality in time)
     - If an item is referenced in memory, it will tend to be referenced again soon.

  b) **Spatial locality** (locality in space)
   - If an item is referenced in memory, nearby items will tend to be referenced soon.

33

33

(a) Temporal Locality

- Recently executed instructions are likely to be executed again very soon.
- *Example*: computing factorial of a number.

```
fact = 1;
for  k = 1 to N
  fact = fact * k;
```

```
        ADDI  $t1,$zero,1
        ADDI  $t2,$zero,N
        ADDI  $t3,$zero,1
Loop:   MUL   $t1,$t1,$t3
        ADDI  $t3,$t3,1
        SGT   $t4,$t3,$t2
        BNEZ  $t4,Loop
```

- The four instructions in the loop are executed more frequently than the others.
  - They exhibit *temporal locality*.

34

34

(b) Spatial Locality

- Instructions residing *close to* a recently executing instruction are likely to be executed soon.

- *Example*: accessing elements of an array.

```
sum = 0;
for  k = 1 to N
   sum = sum + A[k];
```

```
              SUB     $t1,$t1,$t1
              ADDI    $t2,$zero,N
              ADDI    $t3,$zero,1
              ADDI    $t5,$zero,A
      Loop:   LW      $t8,0($t5)
              ADD     $t1,$t1,$t8
              ADDI    $t3,$t3,1
              SGT     $t4,$t3,$t2
              BNEZ    $t4,Loop
```

- Performance can be improved by copying the array into cache memory.

35

35

- An example:
  - Accessing a 2-D array row-wise or column-wise.
  - Assume that the array elements are stored row-wise in memory.
  - One row of data can be brought into cache memory at any given time.
- Which version of the code exhibits more locality?

```
for (i=0; i<128; i++)
  for (j=0; j<128; j++)
    A[i][j] = A[i][j] + 1;
```

```
for (j=0; j<128; j++)
  for (i=0; i<128; i++)
    A[i][j] = A[i][j] + 1;
```

36

18

Performance of Memory Hierarchy

- We first consider a 2-level hierarchy consisting of two levels of memory, say, M_1 and M_2.
  - M_1 can be cache memory, and M_2 can be main memory.

37

---

a) Cost:

- Let c_i denote the cost per bit of memory M_i, and S_i denote the storage capacity in bits of M_i.
- The average cost per bit of the memory hierarchy is given by:

$$\text{Cost } c \ = \ \frac{c_1 S_1 + c_2 S_2}{S_1 + S_2}$$

- In order to have $c \rightarrow c_2$, we must ensure that $S_1 << S_2$.

38

b) Hit Ratio / Hit Rate:

- The hit ratio H is defined as the probability that a logical address generated by the CPU refers to information stored in M_1.
- We can determine H experimentally as follows:
  - A set of representative programs is executed or simulated.
  - The number of references to M_1 and M_2, denoted by N_1 and N_2 respectively, are recorded.

$$H = \frac{N_1}{N_1 + N_2}$$

- The quantity $(1 - H)$ is called the *miss ratio*.

39

c) Access Time:

- Let t_{A1} and t_{A2} denote the access times of M_1 and M_2 respectively, relative to the CPU.
- The average time required by the CPU to access a word in memory can be expressed as:

$$t_A = H.t_{A1} + (1 - H).t_{MISS}$$

where t_{MISS} denotes the time required to handle the miss, called *miss penalty*.

40

- The miss penalty t_{MISS} can be estimated in various ways:

  a) The simplest approach is to set $t_{MISS} = t_{A2}$, that is, when there is a miss the data is accessed directly from M_2.

  b) A request for a word not in M_1 typically causes a block containing the requested word to be transferred from M_2 to M_1. After completion of the block transfer, the word can be accessed in M_1.

  If t_B denotes the block transfer time, we can write

  $$t_{MISS} = t_B + t_{A1} \qquad [\text{since } t_B \gg t_{A1}, \ t_{A2} \approx t_B]$$

  Thus, $t_A = H.t_{A1} + (1 - H).(t_B + t_{A1})$

  c) If t_{HIT} denotes the time required to check whether there is a hit, we can write

  $$t_{MISS} = t_{HIT} + t_B + t_{A1}$$

41

41

d) Efficiency:

- Let $r = t_{A2} / t_{A1}$ denote the access time ratio of the two levels of memory.
- We define the *access efficiency* as $e = t_{A1} / t_A$, which is the factor by which t_A differs from its minimum possible value.

$$\text{Efficiency } e = \frac{t_{A1}}{H.t_{A1} + (1 - H).t_{A2}} = \frac{1}{H + (1 - H).r}$$

42

42

e) Speedup:

- The *speedup* gained by using the memory hierarchy is defined as $S = t_{A2} / t_A$.
- We can write:

$$S = \frac{t_{A2}}{H.t_{A1} + (1-H).t_{A2}} = \frac{1}{(1-H) + H / r}$$

- The same result follows from *Amdahl's law*.

43

43

Some Common Terminologies Used

- **Block**: The smallest unit of information transferred between two levels.
- **Hit Rate**: The fraction of memory accesses found in the upper level.
- **Hit Time**: Time to access the upper level
  - Upper level access time + Time to determine hit/miss
- **Miss**: Data item needs to be retrieved from a block in the lower level.
- **Miss Rate**: The fraction of memory accesses not found in the upper level.
- **Miss Penalty**: Overhead whenever a miss occurs.
  - Time to replace a block in the upper level + Time to transfer the missed block

44

44

Example 1

- Consider a 2-level memory hierarchy consisting of a cache memory M_1 and the main memory M_2. Suppose that the cache is 6 times faster than the main memory, and the cache can be used 90% of the time. How much speedup do we gain by using the cache?

  Here, $r = 6$ and $H = 0.90$

  Thus, $S = 1 / [H / r + (1 - H)] = 1 / (0.90 / 6 + 0.10) = 1 / 0.25 = 4$

45

45

Example 2

- Consider a 2-level memory hierarchy with separate instruction and data caches in level 1, and main memory in level 2.

  The following parameters are given:
  - The clock cycle time is 2 ns.
  - The miss penalty is 15 clock cycles (for both read and write).
  - 1 % of instructions are not found in I-cache.
  - 8 % of data references are not found in D-cache.
  - 20 % of the total memory accesses are for data.
  - Cache access time (including hit detection) is 1 clock cycle.

  Determine the overall average access time.

CPU — I-Cache — Main Memory
CPU — D-Cache — Main Memory

46

46

All memory access

0.80 / Instruction access \ 0.20 Data access

Instruction access: 0.99 / I-cache \ 0.01 Miss

Data access: 0.92 / D-cache \ 0.08 Miss

t_{MISS} = 1 + 15 = 16 cycles

Average number of cycles per access:

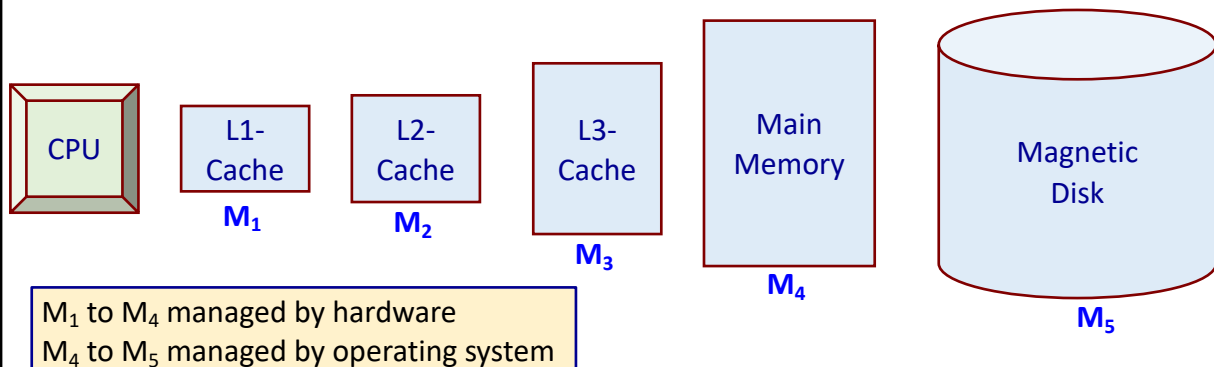   0.80 x (0.99 x 1 + 0.01 x 16) + 0.20 x (0.92 x 1 + 0.08 x 16)

   = 0.92 + 0.44 = 1.36

Thus, average access time t_A = 1.36 x 2 ns = 2.72 ns

47

Performance Calculation for Multi-Level Hierarchy

- Most of the practical memory systems use more than 2 levels of hierarchy.

CPU   L1-Cache M_1 L2-Cache M_2 L3-Cache M_3 Main Memory M_4 Magnetic Disk M_5

M_1 to M_4 managed by hardware
M_4 to M_5 managed by operating system

48

t_{L1} : access time of M_1

t_{L2} : access time of M_2

H_{L1} : hit ratio of M_1

H_{L2} : hit ratio of M_2 with respect to the residual accesses that try to access M_2

- Consider a 3-level hierarchy consisting of L1-cache, L2-cache and main memory.
- Whenever there is a miss in L1, we go to L2.
- Average access time can be calculated as:

$$t_A = H_{L1}.t_{L1} + (1 - H_{L1}).[H_{L2}.t_{L2} + (1 - H_{L2}).t_{MISS2}]$$

- Here, t_{MISS2} is the miss penalty when the requested data is found neither in M_1 nor in M_2.

49

49

Implications of a Memory Hierarchy to the CPU

- Processors designed without memory hierarchy are simpler because all memory accesses take the same amount of time.
  - Misses in a memory hierarchy implies *variable memory access times* as seen by the CPU.

- Some mechanism is required to determine whether or not the requested information is present in the top level of the memory hierarchy.
  - Check happens on every memory access and affects hit time.
  - Implemented in *hardware* to provide acceptable performance.

50

50

- Some mechanism is required to transfer blocks between consecutive levels.
  - If the block transfer requires 10's of clock cycles (like in cache / main memory hierarchy), it is controlled by *hardware*.
  - If the block transfer requires 1000's of clock cycles (like in main memory / secondary memory hierarchy), it can be controlled by *software*.
- Four main questions:
  1. *Block Placement*: Where to place a block in the upper level?
  2. *Block Identification*: How is a block found if present in the upper level?
  3. *Block Replacement*: Which block is to be replaced on a miss?
  4. *Write Strategy*: What happens on a write?

51

51

Cache Memory

52

Introduction

- Let us consider a single-level cache, and that part of the memory hierarchy consisting of cache memory and main memory.

53

53

- Cache memory is logically divided into *blocks* or *lines*, where every block (line) typically contains 8 to 256 bytes.

- When the CPU wants to access a word in memory, a special hardware first checks whether it is present in cache memory.
  - If so (called *cache hit*), the word is directly accessed from the cache memory.
  - If not, the block containing the requested word is brought from main memory to cache.
  - For writes, sometimes the CPU can also directly write to main memory.

- Objective is to keep the commonly used blocks in the cache memory.
  - Will result in significantly improved performance due to the property of *locality of reference*.

54

54

Q1. Where can a block be placed in the cache?

- This is determined by some *mapping algorithms*.
  - Specifies which main memory blocks can reside in which cache memory blocks.
  - At any given time, only a small subset of the main memory blocks can be held in the cache.
- Three common block mapping techniques are used:
  a) **Direct Mapping**
  b) **Associative Mapping**
  c) **(N-way) Set Associative Mapping**

55

An Example Scenario: A 2-level memory hierarchy

- Consider a 2-level cache memory / main memory hierarchy.
  - The cache memory consists of 256 blocks (lines) of 32 words each.
    - Total cache size is 8192 (8K) words.
  - Main memory is addressable by a 24-bit address.
    - Main memory is word addressable.
    - Total size of the main memory is 2^{24} = 16 M words.
    - Number of 32-word blocks in main memory = 16 M / 32 = 512K

56

(a) Direct Mapping

- Each main memory block can be placed in only one block in the cache.
- The mapping function is:

  Cache Block = (Main Memory Block) % (Number of cache blocks)

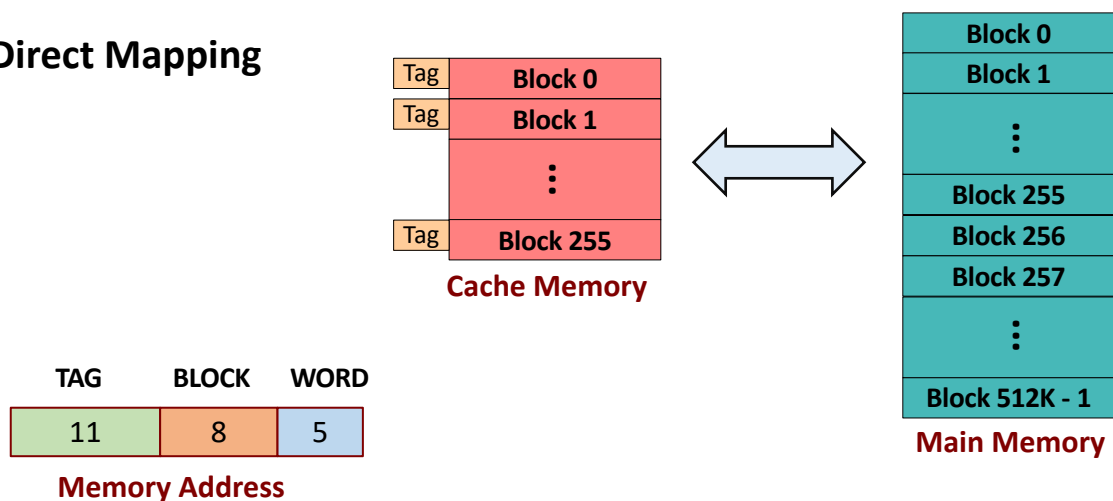- For the example,

  Cache Block = (Main Memory Block) % 256

- Some example mappings:

  0 → 0, 1 → 1, 255 → 255, 256 → 0, 257 → 1, 512 → 0, 512 → 1, etc.

57

57

Direct Mapping

	Cache Memory
Tag	Block 0
Tag	Block 1
	⋮
Tag	Block 255

Main Memory

Block 0
Block 1
⋮
Block 255
Block 256
Block 257
⋮
Block 512K - 1

| TAG	BLOCK	WORD
11	8	5

Memory Address

58

58

- Block replacement algorithm is trivial, as there is no choice.

- More than one MM block is mapped onto the same cache block.
  - May lead to contention even if the cache is not full.
  - New block will replace the old block.
  - May lead to poor performance if both the blocks are frequently used.

- The MM address is divided into three fields: *TAG*, *BLOCK* and *WORD*.
  - When a new block is loaded into the cache, the 8-bit *BLOCK* field determines the cache block where it is to be stored.
  - The high-order 11 bits are stored in a *TAG* register associated with the cache block.
  - When accessing a memory word, the corresponding *TAG* field is compared.
    - Match implies *HIT*.

TAG	BLOCK	WORD
11	8	5

Memory Address

59

59

Direct Mapped Cache
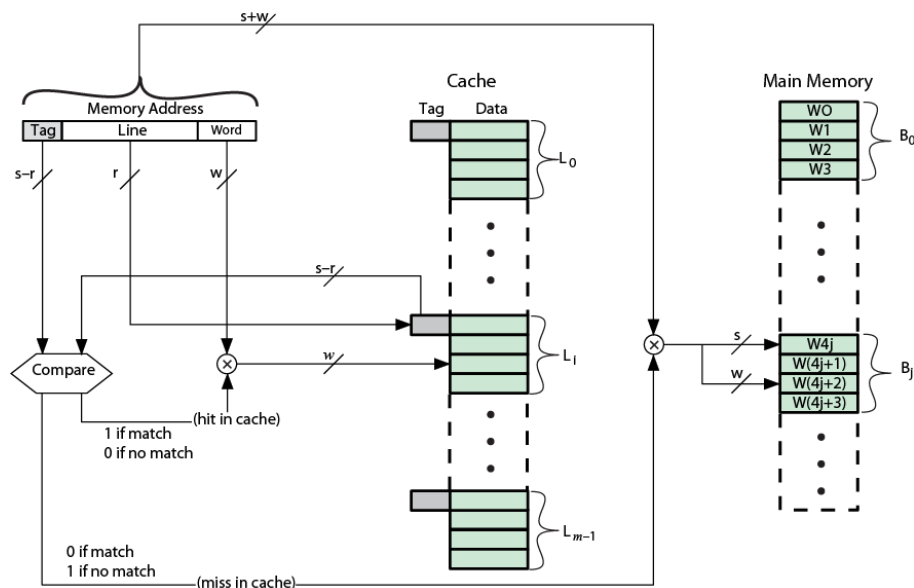
60

(b) Associative Mapping

- Here, a MM block can potentially reside in *any cache block position*.

- The memory address is divided into two fields: *TAG* and *WORD*.
  - When a block is loaded into the cache from MM, the higher order 19 bits of the address are stored into the *TAG* register corresponding to the cache block.
  - When accessing memory, the 19-bit *TAG* field of the address is compared with all the *TAG* registers corresponding to all the cache blocks.

- Requires *associative memory* for storing the *TAG* values.
  - High cost / lack of scalability.

- Because of complete freedom in block positioning, a wide range of replacement algorithms is possible.
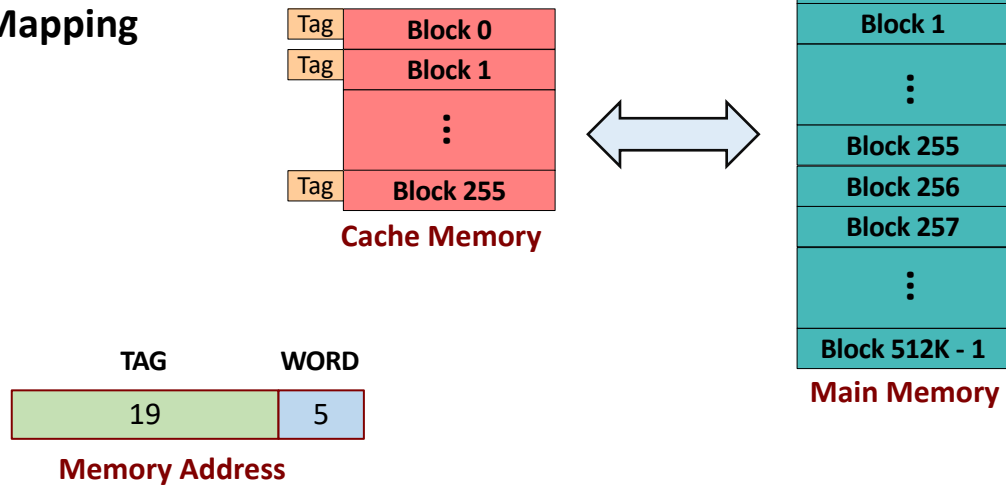
61

61

What is an Associative Memory?

- It is a memory where data is accessed *by contents* rather than address.
  - There is circuitry to compare the applied data value with all the stored values in memory in parallel.
  - Wherever there is a match, the memory system will be returning the information.
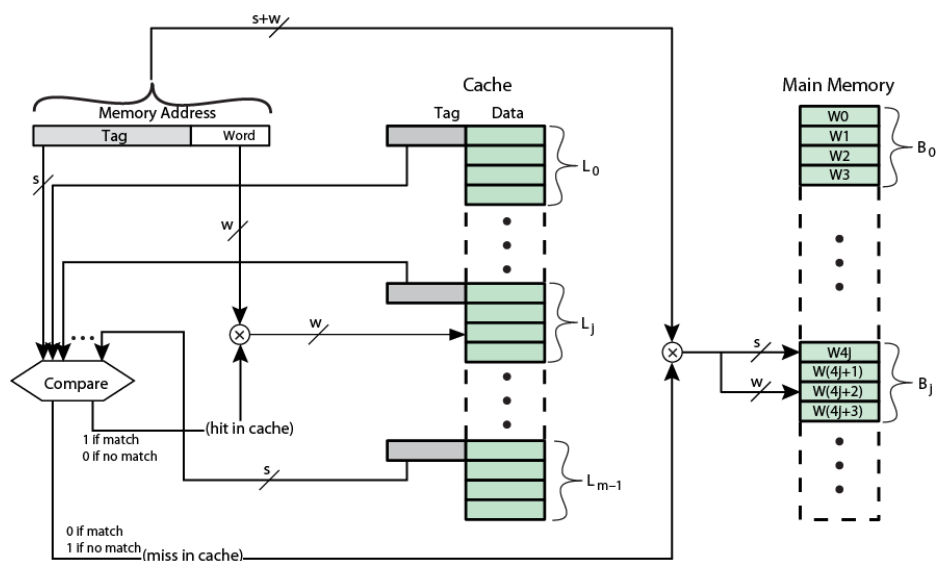 - Very expensive to build – only small capacity units are feasible.

62

Fully Associative Mapping

Tag	Block 0
Tag	Block 1
	⋮
Tag	Block 255

Cache Memory

Block 0
Block 1
⋮
Block 255
Block 256
Block 257
⋮
Block 512K - 1

Main Memory

	TAG	WORD
	19	5

Memory Address

63

Fully Associative Cache Design

(c) N-way Set Associative Mapping

- A group of *N* consecutive blocks in the cache is called a *set*.
- This algorithm is a balance of *direct mapping* and *associative mapping*.
  - Like direct mapping, a MM block is mapped to a set.
    - Set Number = (MM Block Number) % (Number of Sets in Cache)
  - The block can be placed anywhere within the set (there are *N* choices)
- The value of *N* is a design parameter:
  - *N* = 1 :: same as *direct mapping*.
  - *N* = number of cache blocks :: same as *associative mapping*.
 - Typical values of *N* used in practice are: 2, 4 or 8.

65

4-way Set Associative Mapping

TAG	SET	WORD
13	6	5

Memory Address

Set 0
Set 1
⋮
⋮
Set 63
Cache Memory

Block 0
Block 1
⋮
Block 255
Block 256
Block 257
⋮
Block 512K - 1
Main Memory

66

- Illustration for *N = 4*:
  - Number of sets in cache memory = 64.
  - Memory blocks are mapped to a set using *modulo-64 operation*.
  - Example: MM blocks 0, 64, 128, etc. all map to set 0, where they can occupy any of the four available positions.
- MM address is divided into three fields: *TAG*, *SET* and *WORD*.
  - The *TAG* field of the address must be associatively compared to the *TAG* fields of all the 4 blocks of the selected set.
  - This instead of requiring a single large associative memory, we need a number of very small associative memories only one of which will be used at a time.

TAG	SET	WORD
13	6	5

Memory Address

67

67

K-way set-associative cache
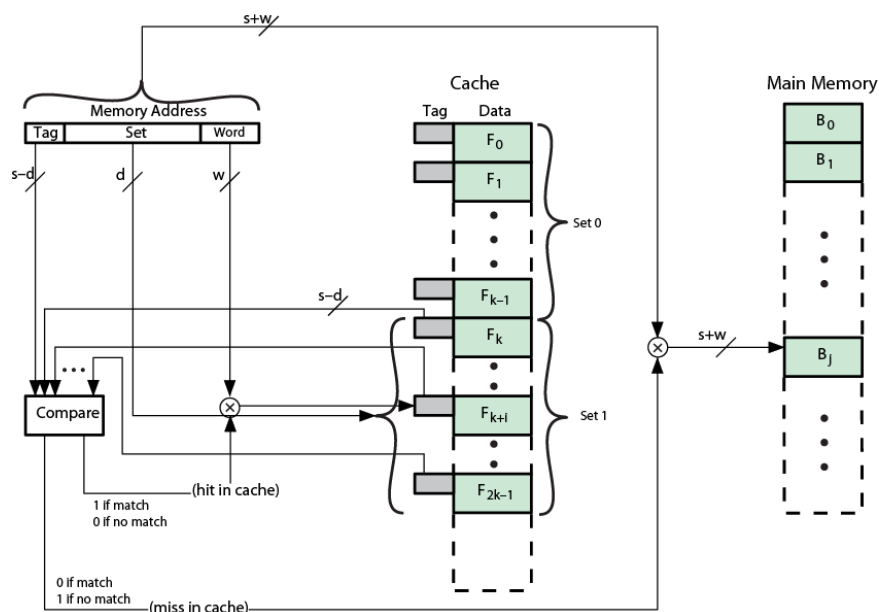
68

Example 3

A 64 KB four-way set-associative cache is byte-addressable and contains 32 B lines.
Memory addresses are 32 b wide.

a) How wide are the tags in this cache?

b) Which main memory addresses are mapped to set number 5?

Solution:

The number of sets in the cache = 64KB/(4 x 32B) = 512.

a. Address (32 b) = 5 b byte offset + 9 b set index + 18 b tag

b. Addresses that have their 9-bit set index equal to 5. These are of the general form $2^{14}a + 2^5 \times 5 + b$; e.g., 160-191, 16 554-16 575, . . .

32-bit address	Tag	Set index	Offset
	18 bits	**9 bits**	**5 bits**

Tag width = 32 – 5 – 9 = 18

Set size = 4×32 B = 128 B
Number of sets = $2^{16}/2^7 = 2^9$

Line width = 32 B = 2^5 B

69

Q2. How is a block found if present in cache?

- Caches include a *TAG* associated with each cache block.
  - The TAG of every cache block where the block being requested may be present needs to be compared with the TAG field of the MM address.
  - All the possible tags are compared in parallel, as speed is important.
- Mapping Algorithms?
  - Direct mapping requires a *single comparison*.
  - Associative mapping requires a *full associative search* over *all the TAGs* corresponding to all cache blocks.
  - Set associative mapping requires a *limited associated search* over the *TAGs of only the selected set*.

70

70

- Use of *valid bit*:
  - There must be a way to know whether a cache block contains *valid* or *garbage* information.
  - A *valid bit* can be added to the *TAG*, which indicates whether the block contains valid data.
  - If the *valid bit is not set*, there is *no need to match* the corresponding *TAG*.

71

71

Q3. Which block should be replaced on a cache miss?

- With fully associative or set associative mapping, there can be several blocks to choose from for replacement when a miss occurs.
- Two primary strategies are used:
  a) *Random*: The candidate block is selected randomly for replacement. This simple strategy tends to spread allocation uniformly.
  b) *Least Recently Used* (LRU): The block replaced is the one that has not been used for the longest period of time.
     - Makes use of a *corollary* of temporal locality:

       "If recently used blocks are likely to be used again, then the best candidate for replacement is the least recently used block"

72

72

- To implement the LRU algorithm, the cache controller must track the LRU block as the computation proceeds.

- Example: Consider a 4-way set associative cache.
  - For tracking the LRU block within a set, we use a 2-bit counter with every block.
  - *When hit occurs*:
    - Counter of the referenced block is reset to 0.
    - Counters with values originally lower than the referenced one are incremented by 1, and all others remain unchanged.
  - *When miss occurs*:
    - If the set is not full, the counter associated with the new block loaded is set to 0, and all other counters are incremented by 1.
    - If the set is full, the block with counter value 3 is removed, the new block put in its place, and the counter set to 0. The other three counters are incremented by 1.

73

73

- It may be verified that the counter values of occupied blocks are all distinct.
- An example:

x Block 0	**x** Block 0	**0** Block 0	**1** Block 0	**2** Block 0	**0** Block 0
x Block 1	**x** Block 1	**x** Block 1	**x** Block 1	**0** Block 1	**1** Block 1
x Block 2	**0** Block 2	**1** Block 2	**2** Block 2	**3** Block 2	**3** Block 2
x Block 3	**x** Block 3	**x** Block 3	**0** Block 3	**1** Block 3	**2** Block 3
Initial	**Miss: Block 2**	**Miss: Block 0**	**Miss: Block 3**	**Miss: Block 1**	**Hit: Block 0**
1 Block 0	**2** Block 0	**2** Block 0	**0** Block 0	**1** Block 0	**1** Block 0
2 Block 1	**3** Block 1	**3** Block 1	**3** Block 1	**0** Block 1	**0** Block 1
0 Block 2	**1** Block 2	**0** Block 2	**1** Block 2	**2** Block 2	**2** Block 2
3 Block 3	**0** Block 3	**1** Block 3	**2** Block 3	**3** Block 3	**3** Block 3
Miss: Block 2	**Hit: Block 3**	**Hit: Block 2**	**Hit: Block 0**	**Miss: Block 1**	**Hit: Block 1**

74

74

Types of Cache Misses

a) Compulsory Miss

- On the first access to a block, the block must be brought into the cache.
- Also known as cold start misses, or first reference misses.
- Can be reduced by increasing cache block size or prefetching cache blocks.

b) Capacity Miss

- Blocks may be replaced from cache because the cache cannot hold all the blocks needed by a program.
- Can be reduced by increasing the total cache size.

75

c) Conflict Miss

- In case of direct mapping or N-way set associative mapping, several blocks may be mapped to the same block or set in the cache.
- May result in block replacements and hence access misses, even though all the cache blocks may not be occupied.
- Can be reduced by increasing the value of N (cache associativity).

76

Q4. What happens on a write?

- Statistical data suggests that read operations (including instruction fetches) dominate processor cache accesses.
  - All instruction fetch operations are read.
  - Most instructions do not write to memory.
- Making the *common case fast*:
  - Optimize cache accesses for reads.
  - But Amdahl's law reminds that for high performance designs we cannot ignore the speed of write operations.

77

77

- The common case (read operations) is relatively easy to make faster.
  - A block(s) can be read at the same time while the *TAG* is being compared with the block address.
  - If the read is a *HIT* the data can be passed to the CPU; if it is a *MISS* ignore it.
- Problems with write operations:
  - The CPU specifies the size of the write (between 1 and 8 bytes), and only that portion of a block has to be changed.
    - Implies a *read-modify-write* sequence of operations on the block.
    - Also, the process of modifying the block cannot begin until the TAG is checked to see if it is a hit.
 - Thus, cache write operations take more time than cache read operations.

78

78

Cache Write Strategies

- Cache designs can be classified based on the write and memory update strategy being used.

 1. *Write Through / Store Through*
  2. *Write Back / Copy Back*

(a) Write Through Strategy

- Information is written to both the cache block and the main memory block.
- Features:
  - Easier to implement.
  - Read misses do not result in writes to the lower level (i.e. MM).
  - The lower level (i.e. MM) has the most updated version of the data – important for I/O operations and multiprocessor systems.
  - A write buffer is often used to reduce CPU write stall time while data is written to main memory.

- **Perfect Write Buffer:**
  - All writes are handled by write buffer; no stalling for write operations.
  - For unified L1 cache,

    Stall Cycles / Memory Access $= $ % Reads $\times$ $(1 - H_{L1}) \cdot t_{MM}$

- **Realistic Write Buffer:**
  - A percentage of write stalls are not eliminated when the write buffer is full.
  - For unified L1 cache,

    Stall Cycles / Memory Access $=$ (% Reads $\times$ $(1 - H_{L1})$ + % write stalls not eliminated) $\times$ t_{MM}

81

(b) Write Back Strategy

- Information is written only to the cache block.
- A modified cache block is written to MM only when it is replaced.
- Features:
  - Writes occur at the speed of cache memory.
  - Multiple writes to a cache block requires only one write to MM.
  - Uses less memory bandwidth, makes it attractive to multiprocessors.
- Write-back cache blocks can be *clean* or *dirty*.
  - A status bit called *dirty bit* or *modified bit* is associated with each cache block, which indicates whether the block was modified in the cache (0: clean, 1: dirty).
  - If the status is clean, the block is not written back to MM while being replaced.

82

**Several writes
to a block**

**Single write
during
replacement**

83

Cache Write Miss Policy

- Since information is usually not needed immediately on a write miss, two options are possible on a cache write miss:

a) **Write Allocate**
  - The missed block is loaded into cache on a write miss, followed by write hit actions.
  - Requires a cache block to be *allocated* for the block to be written into.

b) **No-Write Allocate**
  - The block is modified only in the lower level (i.e. MM), and not loaded into cache.
  - Cache block is *not allocated* for the block to be written into.

84

- Typical usage:

  a) Write-back cache with write-allocate
     - In order to capture subsequent writes to the block in cache.

  b) Write-through cache with no-write-allocate
     - Since subsequent writes still have to to go to MM.

85

85

Estimation of Miss Penalties

- **Write-Through Cache**
  - **Write Hit Operation:**
    - Without write buffer, miss penalty = t_{MM}
    - With perfect write buffer, miss penalty = 0
- **Write-Back Cache**
  - **Write Hit Operation**
    - Miss penalty = 0

86

86

43

- **Write-Back Cache (with Write Allocate)**

  - **Write Hit Operation**
    - Miss penalty = *0*

  - **Read or Write Miss Operation**
    - If the replaced block is clean, miss penalty = t_{MM}
      - No need to write the block back to MM.
      - New block to be brought into MM (t_{MM}).
    - If the replaced block is dirty, miss penalty = 2 t_{MM}
      - Write the block to be replaced to MM (t_{MM}).
      - New block to be brought into MM (t_{MM}).

87

87

Example 4

- Consider a CPU with average CPI of 1.1.
  - Assume an instruction mix: ALU – 50%, LOAD – 15%, STORE – 15%, BRANCH – 20%
  - Assume a cache miss rate of 1.5%, and miss penalty of 50 cycles (= t_{MM}).
  - Calculate the effective CPI for a unified L1 cache, using *write through and no write allocate*, with:
    - a) No write buffer
    - b) Perfect write buffer
    - c) Realistic write buffer that eliminates 85% of write stalls.

  Number of memory accesses per instruction = 1 + 15% + 15% = 1.3
  % Reads = (1 + 0.15) / 1.3 = 88.5%   % Writes = 0.15 / 1.3 = 11.5%

91

91

- **Solution:**

  a) With no write buffer (i.e. stall on all writes)
     - Memory stalls / instr. = 1.3 x 50 x (88.5% x 1.5% + 11.5%) = 8.33 cycles
     - CPI = CPI_{avg} + Memory stalls / instr. = 1.1 + 8.33 = 9.43

  b) With perfect write buffer (i.e. all write stalls are eliminated)
     - Memory stalls / instr. = 1.3 x 50 x (88.5% x 1.5%) = 0.86 cycles
     - CPI = 1.1 + 0.86 = 1.96

  c) With realistic write buffer (85% of write stalls are eliminated)
     - Memory stalls / instr. = 1.3 x 50 x (88.5% x 1.5% + 15% x 11.5%) = 1.98 cycles
     - CPI = 1.1 + 1.98 = 3.08

92

Example 5

- Consider a CPU with average CPI of 1.1.
  - Assume the instruction mix: ALU – 50%, LOAD – 15%, STORE – 15%, BRANCH – 20%
  - Assume a cache miss rate of 1.5%, and miss penalty of 50 cycles (= t_{MM}).
  - Calculate the effective CPI for a unified L1 cache, using *write back and write allocate*, with the probability of a cache block being dirty is 10%.

    Number of memory accesses per instruction = 1 + 15% + 15% = 1.3

93

- **Solution:**
    - Memory accesses per instruction = 1.3
    - Stalls / access = $(1 - H_{L1})$. $(t_{MM}$ x % clean + $2t_{MM}$ x % dirty)
                    = 1.5% x (50 x 90% + 100 x 10%) = 0.825 cycles
    - Average memory access time = 1 + stalls / access = 1 + 0.825 = 1.825 cycles
    - Memory stalls / instr. = 1.3 x 0.825 = 1.07 cycles
    - Thus, effective CPI = 1.1 + 1.07 = 2.17

94

94