

Computer Organization and Architecture

Module 5 (Part 2) **Design of Memory Subsystems**

Prof. Indranil Sengupta

Dr. Sarani Bhattacharya

Department of Computer Science and Engineering

IIT Kharagpur

1

Memory Interfacing and Addressing

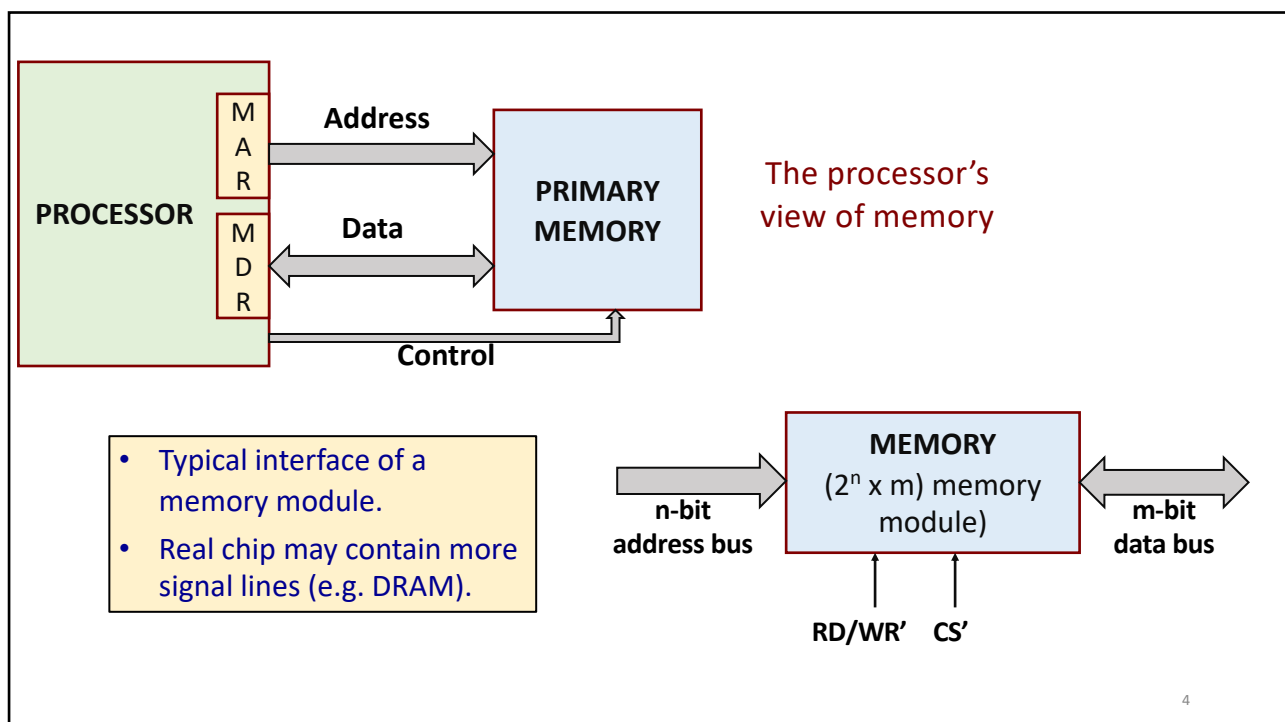
2

Memory Interfacing

- Basic problem:
 - Interfacing one or more memory modules to the processor.
 - We assume a single level memory at present (i.e. no cache memory).
- Questions to be answered:
 - How the processor address and data lines are connected to memory modules?
 - How are the addresses decoded?
 - How are the memory addresses distributed among the memory modules?
 - How to speed up data transfer rate between processor and memory?

3

3



4

4


A Note About the Memory Interface Signals

- The data signals of a memory module (RAM) are typically bidirectional.
 - Some memory chips may have separate data in and data out lines.
- For memory **READ** operation:
 - Address of memory location is applied to **address lines**.
 - **RD/WR'** control signal is set to 1, and **CS'** is set to 0.
 - Data is read out through the **data lines** after memory access time delay.
- For memory **WRITE** operation:
 - Address of memory location is applied to **address lines**, and the data to be written to **data lines**.
 - **RD/WR'** control signal is set to 0, and **CS'** is set to 0.

5

5

- Why is **CS'** signal required?
 - To handle multiple memory modules interfacing problem.
 - We typically select only one out of several memory modules at a time.
- What happens when **CS' = 1**?
 - When a memory module is **not selected**, the data lines are set to the **high impedance state** (i.e. electrically disconnected).
 - An example scenario is shown.



6

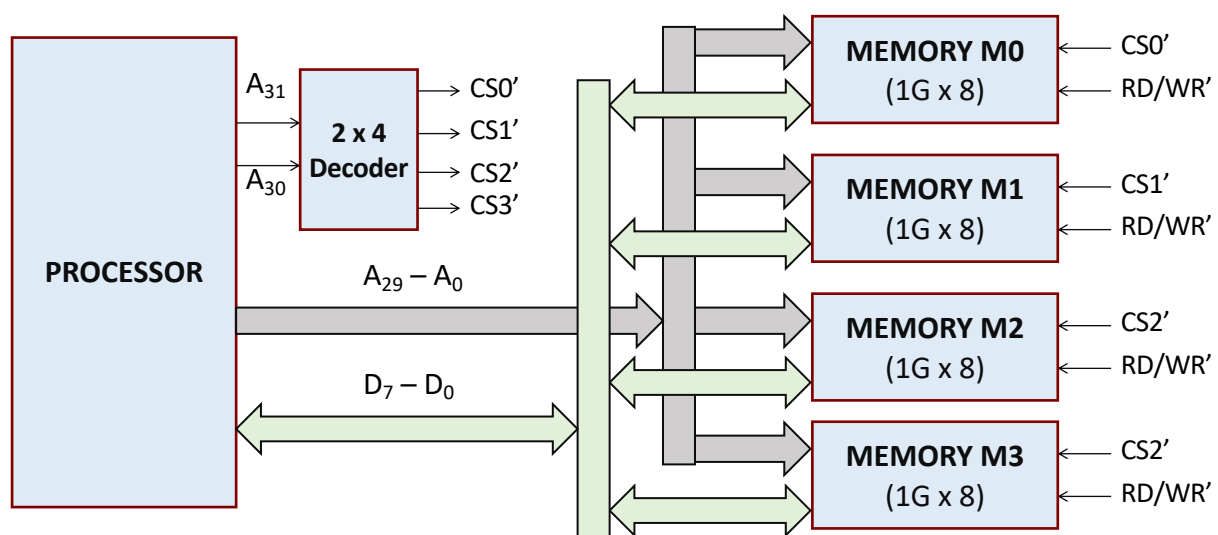
6

An Example Memory Interfacing Problem

- Consider a MIPS32 like processor with a 32-bit address.
 - Maximum memory that can be connected is $2^{32} = 4$ Gbytes.
 - Assume that the processor data lines are 8 bits.
- Assume that memory chips (RAM) are available with *size 1 Gbyte*.
 - 30 address lines and 8 data lines.
 - Low-order 30 address lines ($A_{29}-A_0$) are connected to the memory modules.
- We want to interface *4 such chips* to the processor.
 - Total memory of 4 Gbytes.

7

7



8

8

- High order address lines (A_{31} and A_{30}) select one of the memory modules.

- **When is M0 selected?**

- Address is: **0 0** x
- Range of addresses is: 0x**0**0000000 to 0x**3**FFFFFFF

- **When is M1 selected?**

- Address is: **0 1** x
- Range of addresses is: 0x**4**0000000 to 0x**7**FFFFFFF

- **When is M2 selected?**

- Address is: **1 0** x
- Range of addresses is: 0x**8**0000000 to 0x**B**FFFFFFF

- **When is M3 selected?**

- Address is: **1 1** x
- Range of addresses is: 0x**C**0000000 to 0x**F**FFFFFFF

9

9

- **An observation:**

- Consecutive block of bytes are mapped to the same memory module.
- For MIPS32, we have to access 32 bits (4 bytes) of data in parallel, which requires four sequential memory accesses here.
- We shall look at an alternate memory organization later that would make this possible.
 - Called *memory interleaving*.


10

10

Exercise 1

- 6116 is a 2K x 8 RAM chip. Build a 16K x 8 memory using 6116 chips. Show all the connections, and state how the addresses are distributed across memory modules.

6116 Pinout

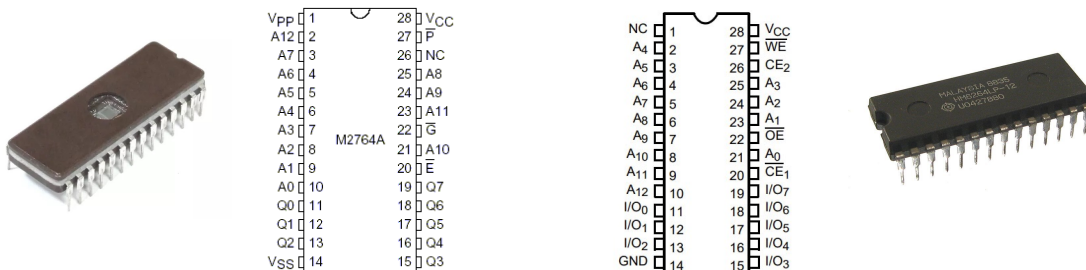


11

Exercise 2

- 6264 is an 8K x 8 RAM chip, and 2764 is an 8K x 8 EPROM chip. Build a 32K x 8 memory system with 8 KB EPROM and 24 KB RAM. Show all the connections, and state how the addresses are distributed across memory modules.

6264 Pinout




12

Exercise 3

- 2114 is a 1K x 4 RAM chip. Build a 4K x 8 memory using 2114 chips. Show all the connections, and state how the addresses are distributed across memory modules.

2114N Pinout



13

Memory Interleaving

14

Improved Memory Interface for MIPS32

- We make small changes in the organization so that 32-bits of data can be fetched in a single memory access cycle.
 - Exploit the concept of *memory interleaving*.
 - Consecutive bytes are mapped to different memory modules.
- The main changes:
 - High order 30 address lines (A_{31} - A_2) are connected to memory modules.
 - Low order two address lines (A_1 and A_0) are used to select one of the modules.

15

15

- How are the addresses mapped to memory modules?
 - **Module M0**: 0, 4, 8, 12, 16, 20, 24, ...
 - **Module M1**: 1, 5, 9, 13, 17, 21, 25, ...
 - **Module M2**: 2, 6, 10, 14, 18, 22, 26, ...
 - **Module M3**: 3, 7, 11, 15, 19, 23, 27, ...
- Memory addresses are *interleaved* across memory modules.
- What we can gain from this mapping?
 - Consecutive addresses are mapped to consecutive modules.
 - Possible to access four consecutive words in the same cycle, if all four modules are enabled simultaneously.

16

16

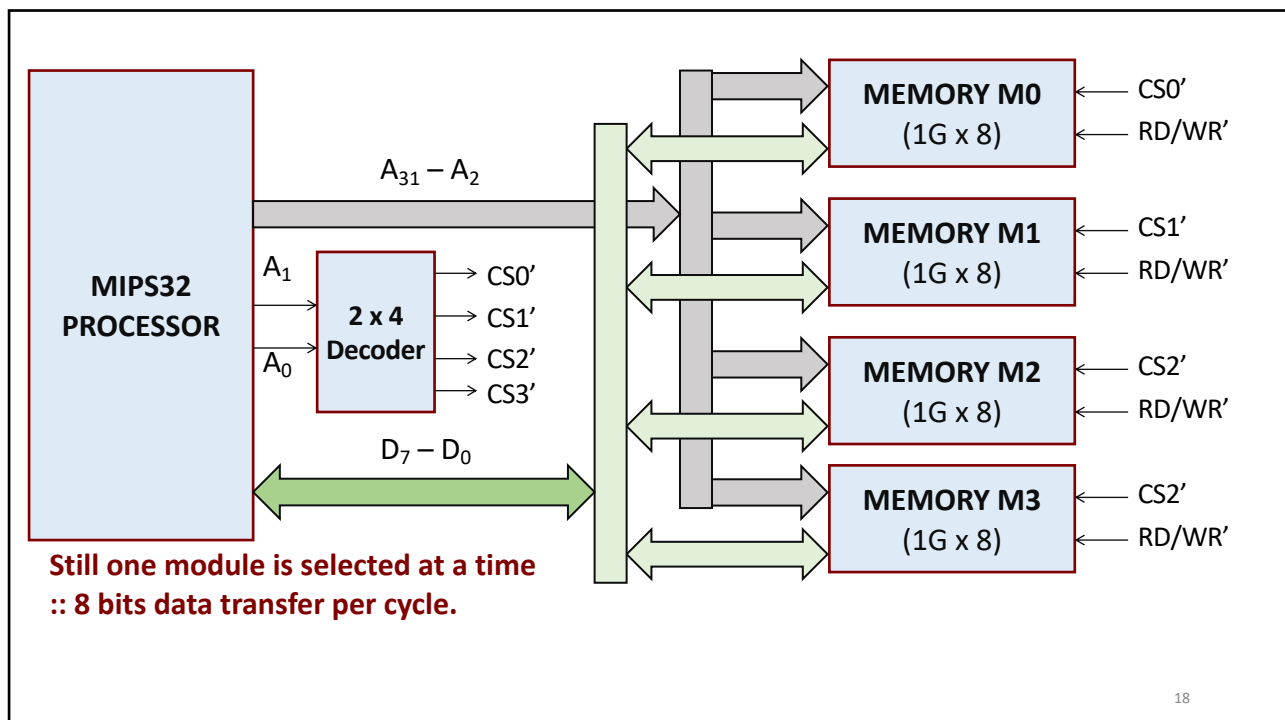
- Motivation for word alignment in MIPS32 data words.
 - 32-bit words start from a memory address that is divisible by 4.
 - Corresponding byte addresses are (0, 1, 2, 3), (4, 5, 6, 7), (8, 9, 10, 11), (12, 13, 14, 15), etc. → *last two bits of the addresses are 00, 01, 10 and 11.*
 - Possible to transfer all the four bytes of a word in a *single memory cycle.*
 - What happens if a word is not aligned?

2 memory cycles required

 - Say: (1, 2, 3, 4) or (2, 3, 4, 5) or (3, 4, 5, 6).
 - Two of the bytes will be mapped to the same memory module.
 - Hence the word cannot be transferred in a single memory cycle.

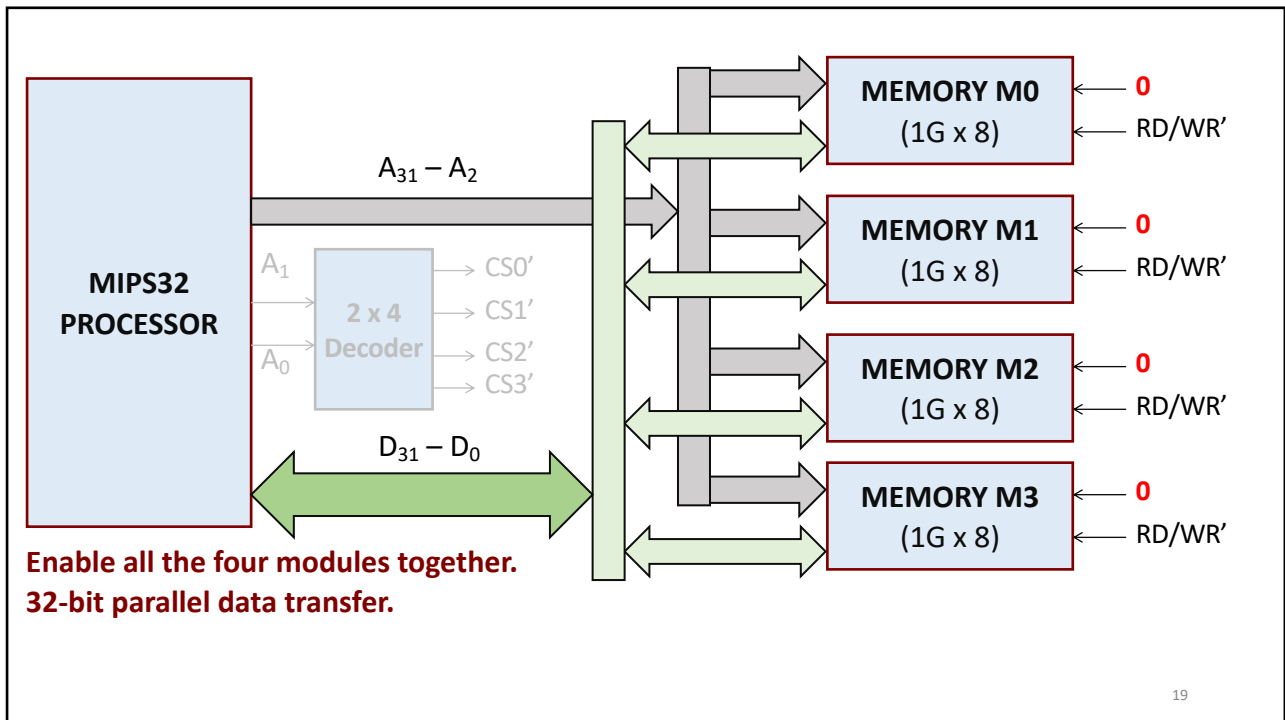
17

17



18

18



19

Memory Latency and Bandwidth

• Memory Latency:

- The delay from the issue of a memory read request to the first byte of data becoming available.

• Memory Bandwidth:

- The maximum number of bytes that can be transferred between the processor and the memory system per unit time.

20

20

- **Example 1:**

Consider a memory system that takes 20 ns to service the access of a single 32-bit word.

Latency $L = 20$ ns per 32-bit word.

Bandwidth $BW = 32 / (20 \times 10^{-9}) = 200$ Mbytes per second.

- **Example 2:**

The memory system is modified to accept a new (still 20ns) request for a 32-bit word every 5 ns by overlapping requests.

Latency $L = 20$ ns per 32-bit word (*no change*).

Bandwidth $BW = 32 / (5 \times 10^{-9}) = 800$ Mbytes per second.

21