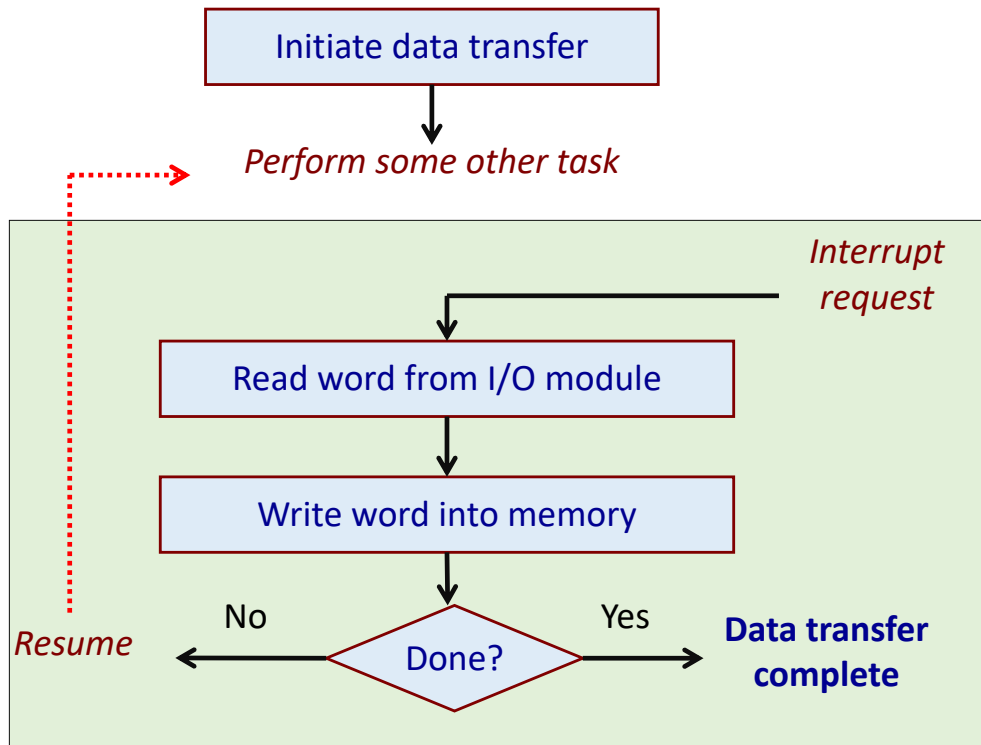


(c) Interrupt-Driven Data Transfer

- The CPU initiates the data transfer and proceeds to perform some other task.
- When the I/O module is ready for data transfer, it informs the CPU by activating a signal (called *interrupt request*).
- The CPU suspends the task it was doing, services the request (that is, carries out the data transfer), and returns back to the task it was doing.
- Characteristics:
 - CPU time is not wasted while checking the status of the I/O module.
 - CPU time is required only during data transfer, plus some overheads for transferring and returning control.




This part of the program that gets activated when interrupt request comes is called *interrupt handler* or *interrupt service subroutine (ISS)*.

31

Some Features of Interrupt-Driven Data Transfer

- How is ISS different from a normal subroutine or function?
 - A function is called from well-defined places in the calling program.
 - Only the relevant registers need to be saved on entry to the function, and restored before return.
 - The ISS can get invoked from anywhere in the program that was executing.
 - Depends on when the interrupt request signal arrived.
 - So potentially all the registers that are used in the ISS needs to be saved and restored.



- We shall learn later why the *INTA* signal is required.

Some Challenges in Interrupts

- For multiple sources of interrupts, how to know the address of the ISS?
- How to handle multiple interrupts?
 - While an interrupt request is being processed, another interrupt request might come.
 - Enabling, disabling and masking of interrupts.
- How to handle simultaneously arriving interrupts?
- Sources of interrupts other than I/O devices.
 - Exceptions, TRAP / System Call, etc.


What happens when an interrupt request arrives?

- At the end of the current instruction execution, the *PC* and *program status word* (PSW) are saved in stack automatically.
 - PSW contains status flags and other processor status information.
- The interrupt is acknowledged, the interrupt vector obtained, based on which control transfers to the appropriate ISS.
 - Different interrupting devices may have different ISS's.
- After handling the interrupt, the ISR executes a special *Return From Interrupt* (RTI) instruction.
 - Restores the PSW and returns control to the saved PC address.
 - Unlike normal RETURN where PSW is not restored.

35

- An instruction cycle typically consists of several machine cycles.
 - For MIPS32, there are 5 machine cycles.

*Interrupt not
acknowledged
here*

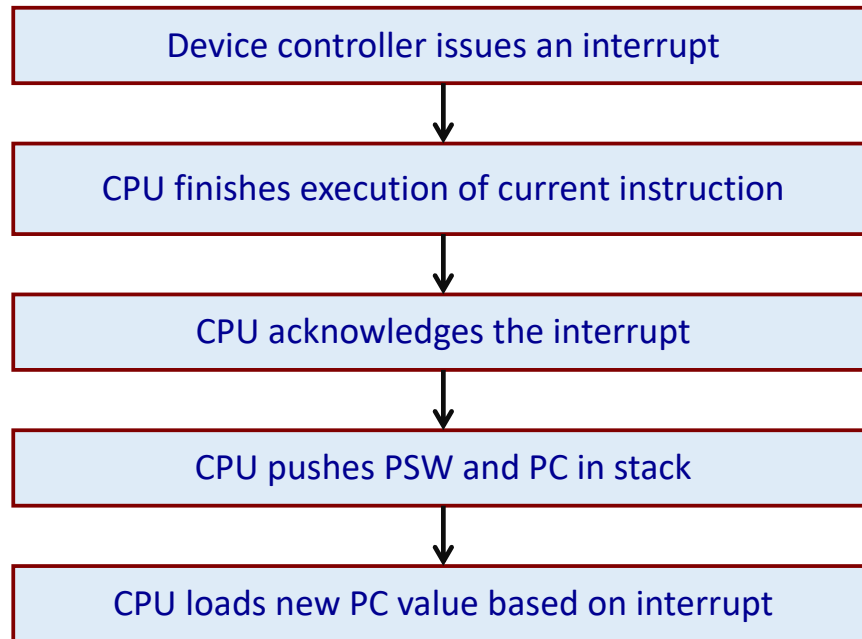


*Interrupt
acknowledged
here*

← Instruction Cycle →

General Interrupt Processing

By hardware



Save remainder of process state information

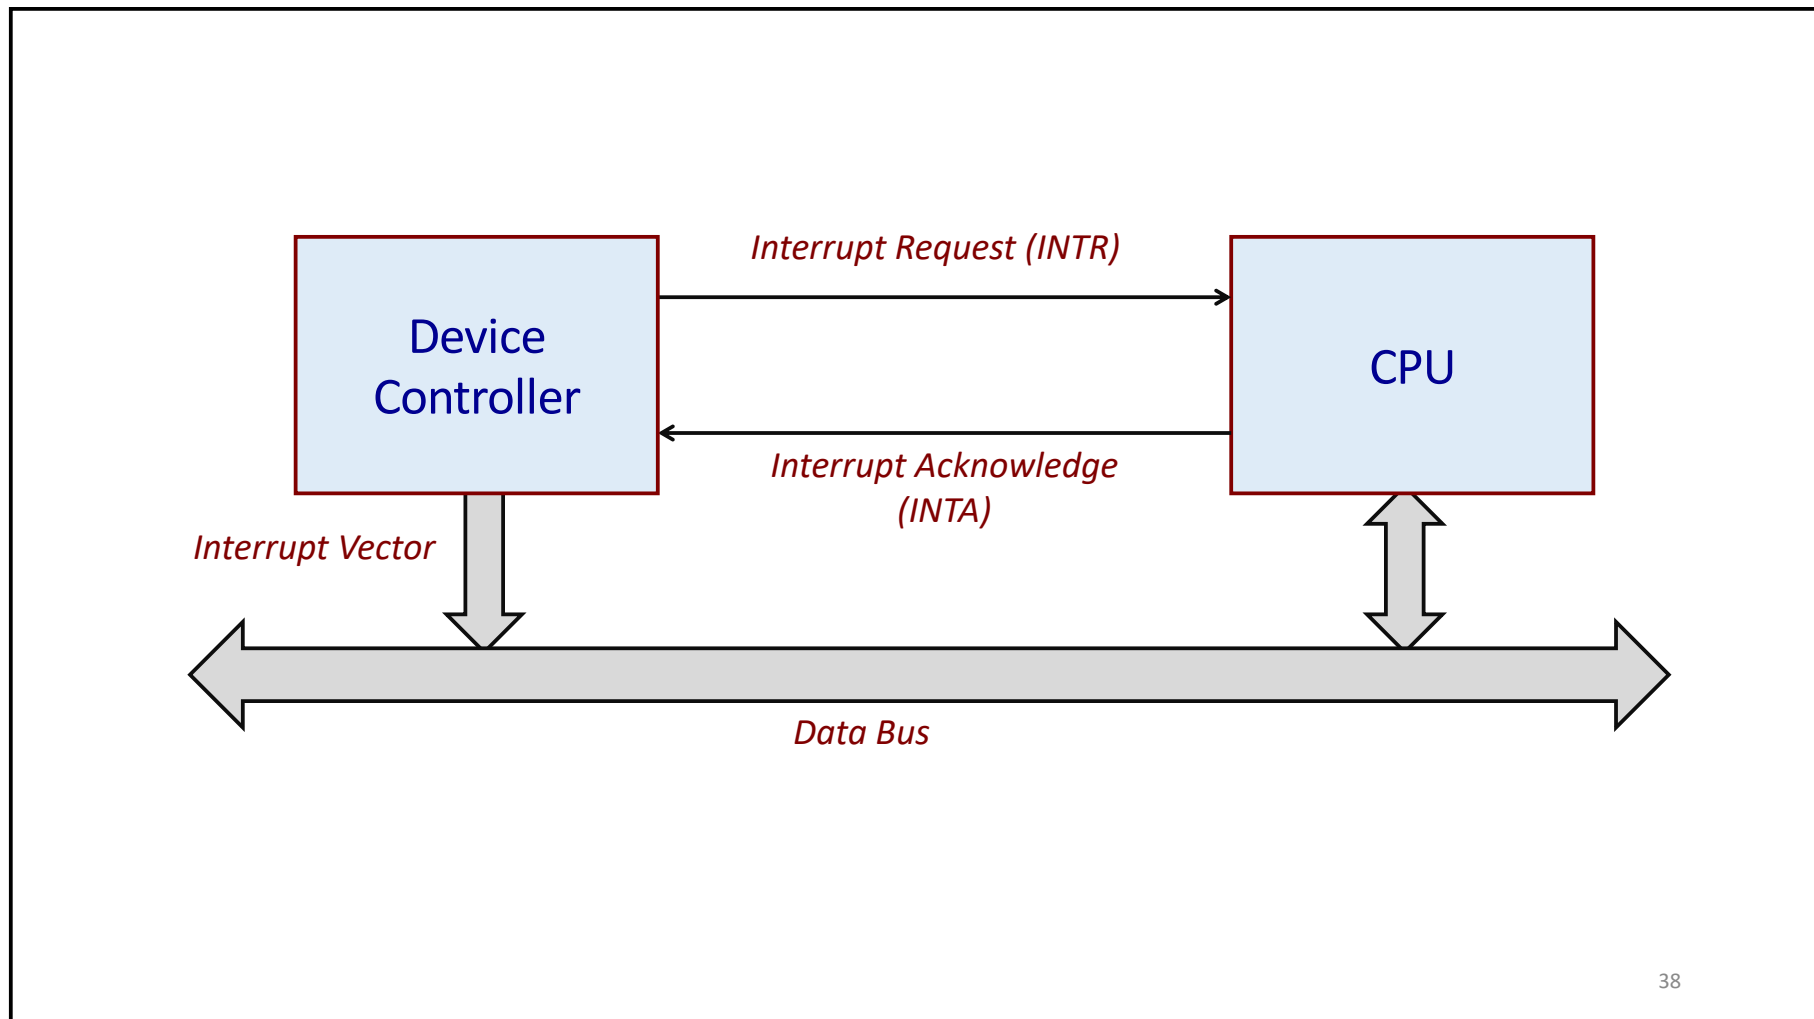
Process interrupt request

Restore process state information

Restore old PSW and PC

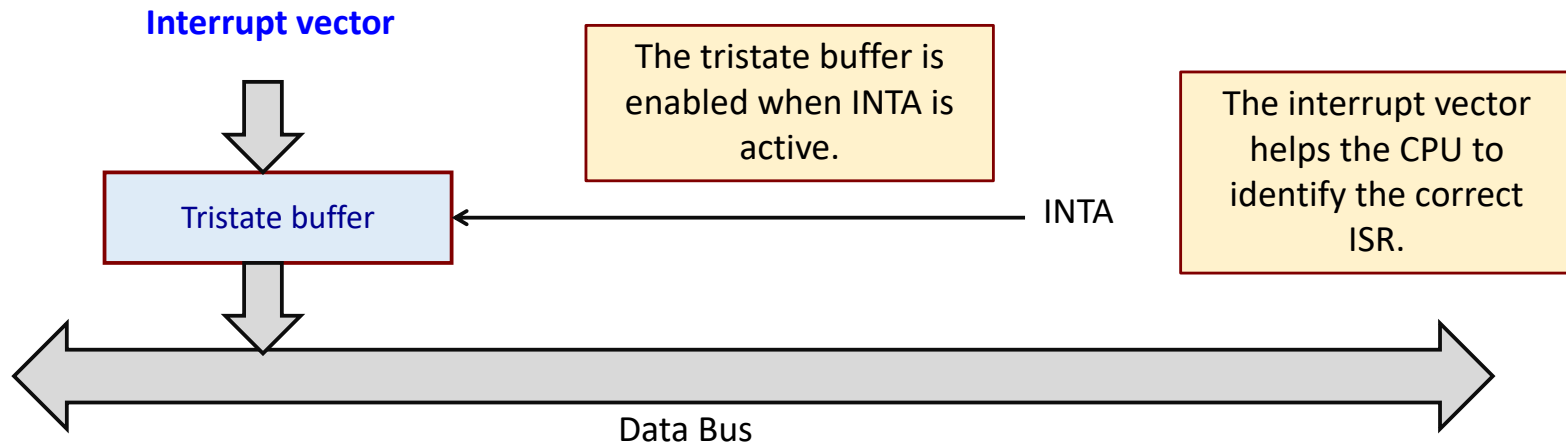
By software

37



- How is the interrupt vector sent on the data bus in response to INTA?

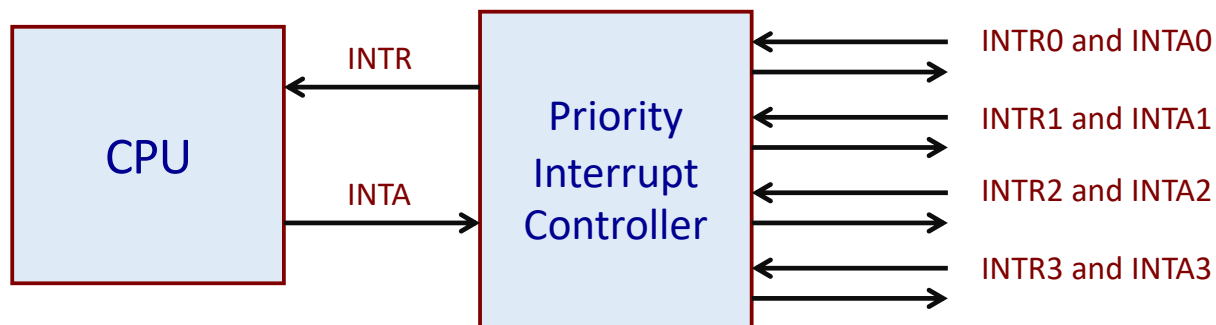
- a) Device controller sends *INTR* to the CPU.
- b) CPU finishes the current instruction and sends back *INTA*.
- c) Device controller sends *interrupt vector* (or number) over data bus.
- d) CPU reads the interrupt vector, and identifies the device.



39

Multiple Devices Interrupting the CPU

- A common solution is to use a *priority interrupt controller*.
 - The interrupt controller interacts with CPU on one side and multiple devices on the other side.
 - For simultaneous interrupt requests, interrupt priority is defined.
 - The interrupt controller is responsible for sending the interrupt vector to CPU.



- How it works?

- The *INTR* line is made active when one or more of the device(s) activate their interrupt request line.

$$INTR = INTRO + INTR1 + INTR2 + INTR3$$

- When the CPU sends back *INTA*, the interrupt controller sends back the corresponding acknowledge to the interrupting device, and puts the interrupt vector on the data bus.
- The interrupt controller is programmable, where the interrupt vectors for the various interrupts can be programmed (specified).
- For more than one interrupt request simultaneously active, a priority mechanism is used (e.g. *INTRO* is highest priority, followed by *INTR1*, etc.).

41

How is interrupt nesting handled?

- Consider the scenario:
 - a) A device **D0** has interrupted and the CPU is executing the ISS for **D0**.
 - b) In the mean time, another device **D1** has interrupted.
- Two possible scenarios here:
 - 1) **D1** will interrupt the ISS for **D0**, get processed first, and then the ISS for **D0** will be resumed. → **CREATES PROBLEM FOR MULTI NESTING**
 - 2) Disable the interrupt system automatically whenever an interrupt is acknowledged so that handling of nested interrupts is not required.

- Typical instruction set architectures have the following instructions:
 - EI : Enable interrupt
 - DI : Disable interrupt
- For the second scenario as discussed, the ISS will give an *EI* instruction just before *RTI*.
 - Some ISA combine *EI* and *RTI* in a single instruction.
- The *DI* instruction is sometimes used by the operating system to execute atomic code (e.g. semaphore wait and signal operations).
 - Nobody should interrupt the code while it is being executed.

43

Cases that make interrupt handling difficult

- For some interrupts, it is not possible to finish the execution of the current instruction.
 - A special *RETURN* instruction is required that would return and restart the interrupted instructions.
- Some examples:
 - a) Page fault interrupt:** A memory location is being accessed that is not presently available in main memory.
 - b) Arithmetic exception:** Some error has occurred during some arithmetic operation (e.g. division by zero).

Handling Multiple Devices

- Suppose that a number of devices capable of generating interrupts are connected to the CPU.
- The following questions need to be answered.
 - a) How can the CPU identify the interrupting device?
 - b) How can the CPU obtain the starting address of the appropriate ISS?
 - c) Should interrupt nesting be allowed?
 - d) How should two or more simultaneous interrupt requests be handled?

45

(a) Device Identification

- Suppose that an external device requests an interrupt by activating an *INTR* line that is common to all the devices. That is,

$$INTR = INTR_1 + INTR_2 + \dots + INTR_n$$

- Each device can have a status bit indicating whether it has interrupted.
 - CPU can *poll* the status bits to find out who has interrupted.
- A better alternative is to use the *interrupt vector* concept discussed earlier.
 - The interrupting device sends a special identifying code on the data bus upon receiving the interrupt acknowledge.

(b) Find Starting Address of ISS

- For a processor with multiple interrupt request inputs, the address of the ISS can be fixed for each individual input.
 - Lacks flexibility.
- If we use the interrupt vector scheme discussed earlier, the device is able to identify itself to the CPU.
 - CPU can then lookup a table where the ISS addresses for all the devices are stored.
 - The interrupt latency is somewhat increased, since we are not immediately jumping to the ISS.

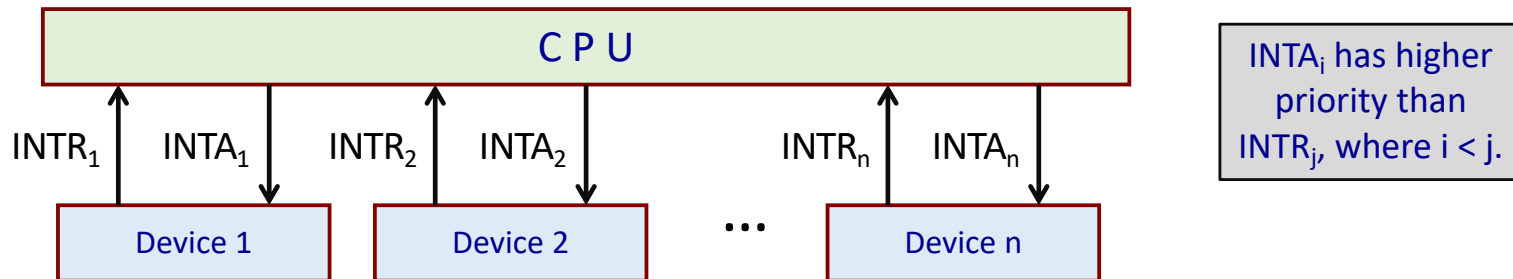
47

(c) Interrupt Nesting

- A simple approach:
 - Disable all interrupts during the execution of an ISS.
 - This ensures that the interrupt request from one device will not cause more than one interruptions.
 - ISS's are typically short, and the delay they may cause in handling a second interrupt request is often acceptable.
- *Interrupt priority:*
 - Some interrupting devices may be assigned higher priorities than others.
 - Example: timer interrupt to maintain a real-time clock.
 - Higher priority interrupt may interrupt the ISS of lower priority ones.

(d) Simultaneous Requests

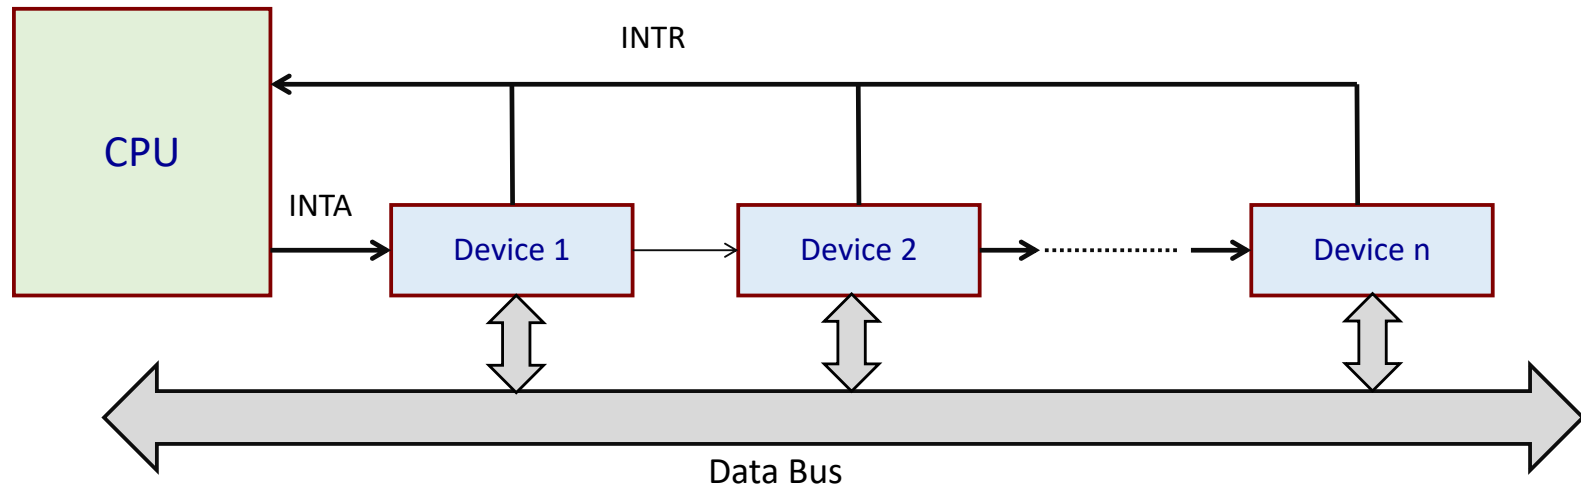
- Here we consider the problem of simultaneous arrivals of interrupt requests from two or more devices.
 - CPU should have some mechanism by which only one request is serviced while the others are delayed or ignored.
 - If the CPU has multiple interrupt request lines, it can have a *priority scheme* where it accepts the request with the highest priority.



49

- Another way to assign priority is to use polling using *daisy chaining*.
 - In polling, priority is automatically assigned based on the order in which the devices are polled.
 - In daisy chain connection, the *INTR* line is common to all the devices, but the *INTA* line is connected in a daisy chain fashion allowing it to propagate serially through the devices.
 - A device when it receives *INTA*, passes the signal to the next device only if it had not interrupted. Else, it stops the propagation of *INTA*, and puts the identifying code on the data bus.
 - Thus, the device that is electrically closest to the CPU will have the highest priority.

Daisy Chain Arrangement

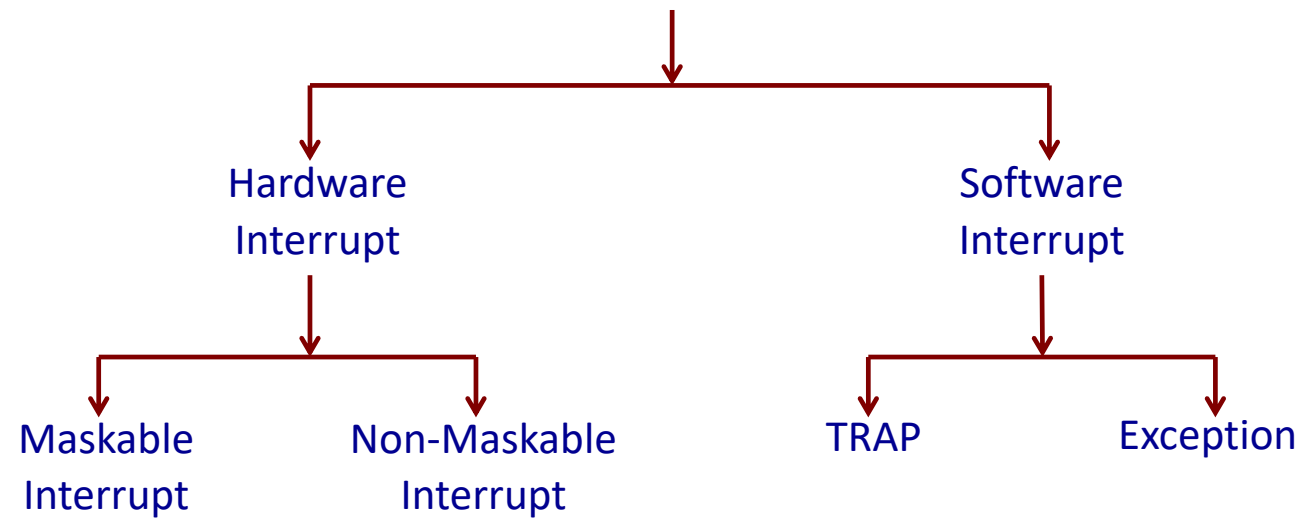


51

51

51

Types of Interrupts



- *Hardware Interrupt:*

- The interrupt signal is coming from a device external to the CPU.
- Example: keyboard interrupt, timer interrupt, etc.

- *Maskable Interrupt:*

- Hardware interrupts that can be masked or delayed when a higher priority interrupt request arrives.
- There are processor instructions that can selectively mask and unmask the interrupt request lines of the CPU.

53

- *Non-Maskable Interrupt:*
 - Interrupts that cannot be delayed and should be handled by the CPU immediately.
 - Examples: power fail interrupts, real-time system interrupts, etc.
- *Software Interrupt:*
 - They are caused due to execution of some instructions.
 - Not caused due to external inputs.
- *TRAP:*
 - They are special instructions used to request services from the operating system.
 - Also called *system calls*.
- *Exception:*
 - These are unplanned interrupts generated while executing a program.
 - They are generated from within the system.
 - Examples: invalid opcode, divide by zero, page fault, invalid memory access, etc.

Direct Memory Access (DMA)

Introduction

- In the data transfer methods discussed under programmed I/O, it is assumed that machine instructions are used to transfer the data between I/O device and memory.
 - Not very suitable when large blocks of data are required to be transferred at high speed (e.g. transfer of a disk block).
- An alternate approach is *Direct Memory Access* (DMA).
 - Allows transfer of a block of data directly between an I/O device and memory, without continuous CPU intervention.

- Why programmed I/O is not suitable for high-speed data transfer?
 - a) Several program instructions have to be executed for each data word transferred between the I/O device and memory.
 - Suppose 20 instructions are required for each word transfer.
 - The CPI of the machine running at 1 GHz clock is 1.
 - So, 20 nsec is required for each word transfer → maximum 50 M words/sec
 - Data transfer rates of fast disks are higher than this figure.

57

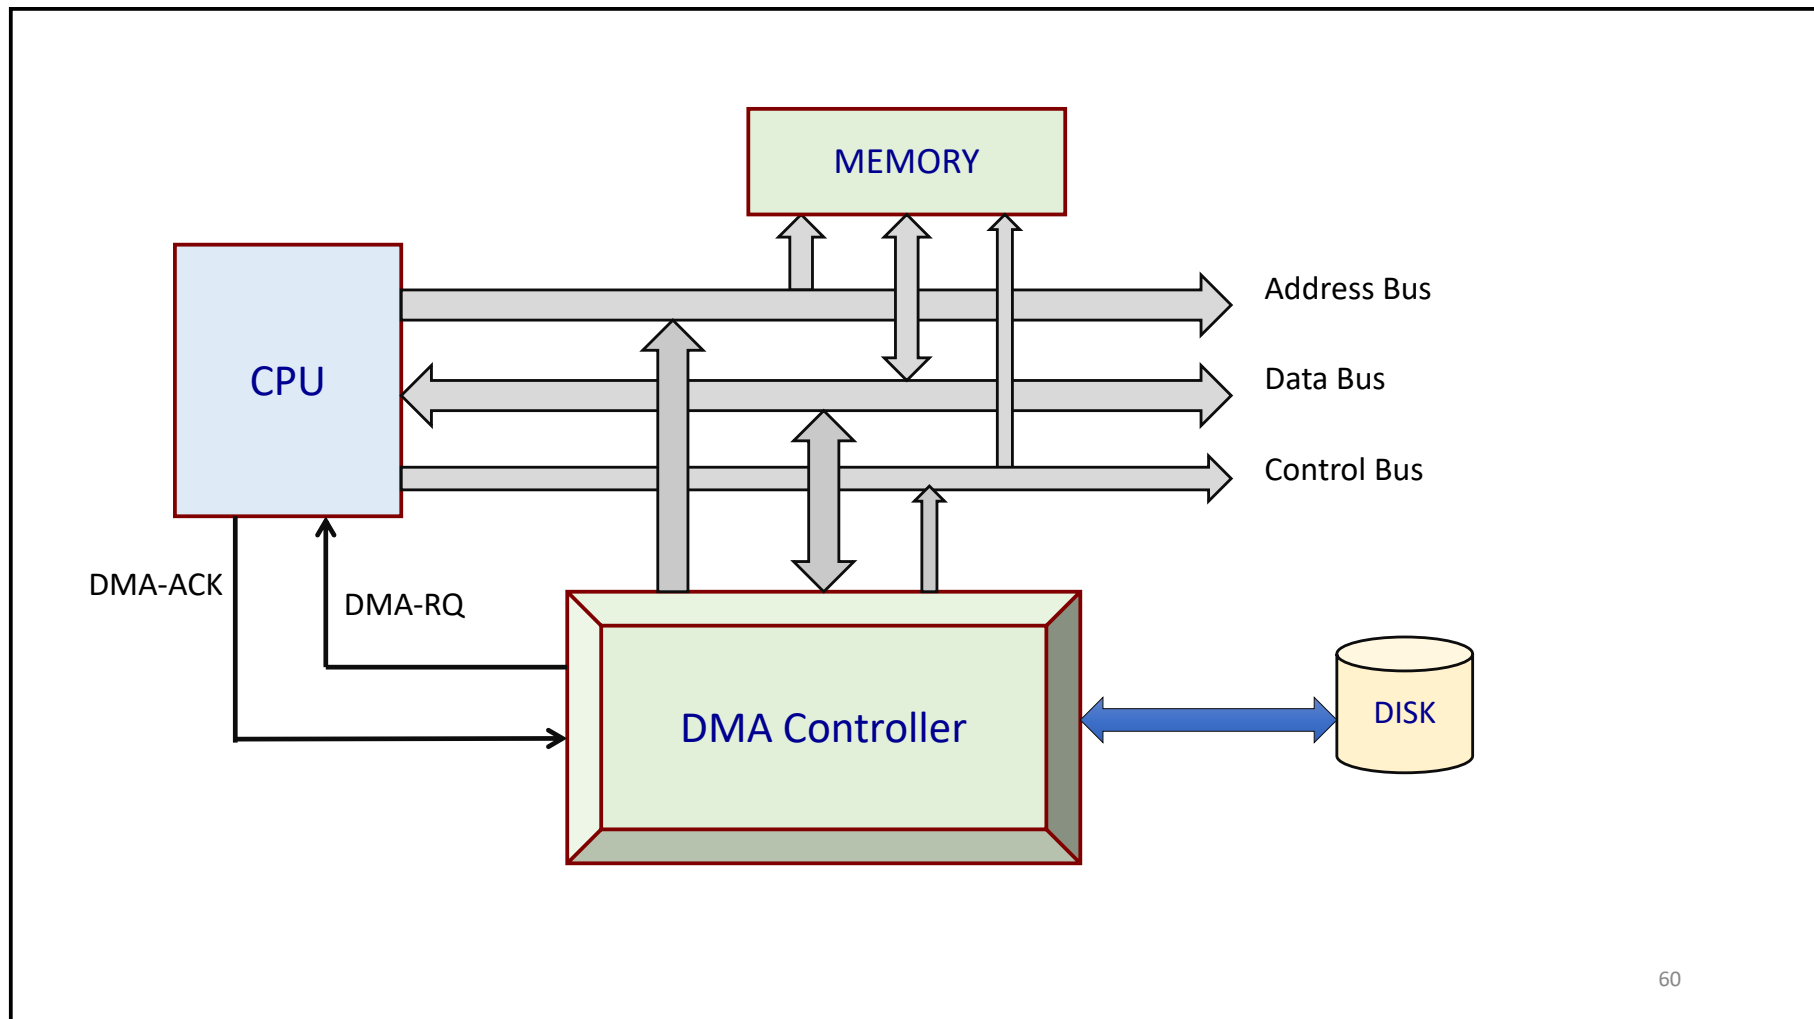
b) Many high speed peripheral devices like disk have a synchronous mode of operation, where data are transferred at a fixed rate.

- Consider a disk rotating at 7200 rpm, with average rotational delay of 4.15 msec.
- Suppose there are 64 Kbytes of data recorded in every track.
- Once the disk head reaches the desired track, there will be a sustained data transfer at rate $64 \text{ Kbytes} / 4.15 \text{ msec} = 15.4 \text{ MBps}$.
- This sustained data transfer rate is comparable to the memory bandwidth, and cannot be handled by programmed I/O.

DMA Controller

- A hardwired controller called the *DMA controller* can enable direct data transfer between I/O device (e.g. disk) and memory without CPU intervention.
 - No need to execute instructions to carry out data transfer.
 - Maximum data transfer speed will be determined by the rate with which memory read and write operations can be carried out.
 - Much faster than programmed I/O.

59



Steps Involved

- a) When the CPU wants to transfer data, it initializes the DMA controller.
 - How many bytes to transfer, address in memory for the transfer.
- b) When the I/O device is ready for the transfer, the DMA controller sends *DMA-RQ* signal to the CPU.
- c) CPU waits till the next DMA breakpoint, relinquishes control of the bus (i.e. puts them in high impedance state), and sends *DMA-ACK* to DMA controller.
- d) Now DMA controller enables its bus interface, and transfers data directly to/from memory.
- e) When done, it deactivates the *DMA-RQ* signal.
- f) The CPU again begins to use the bus to access memory.

61

- The DMA breakpoints:
 - DMA request can be acknowledged at the end of any machine cycle.

DMA breakpoints

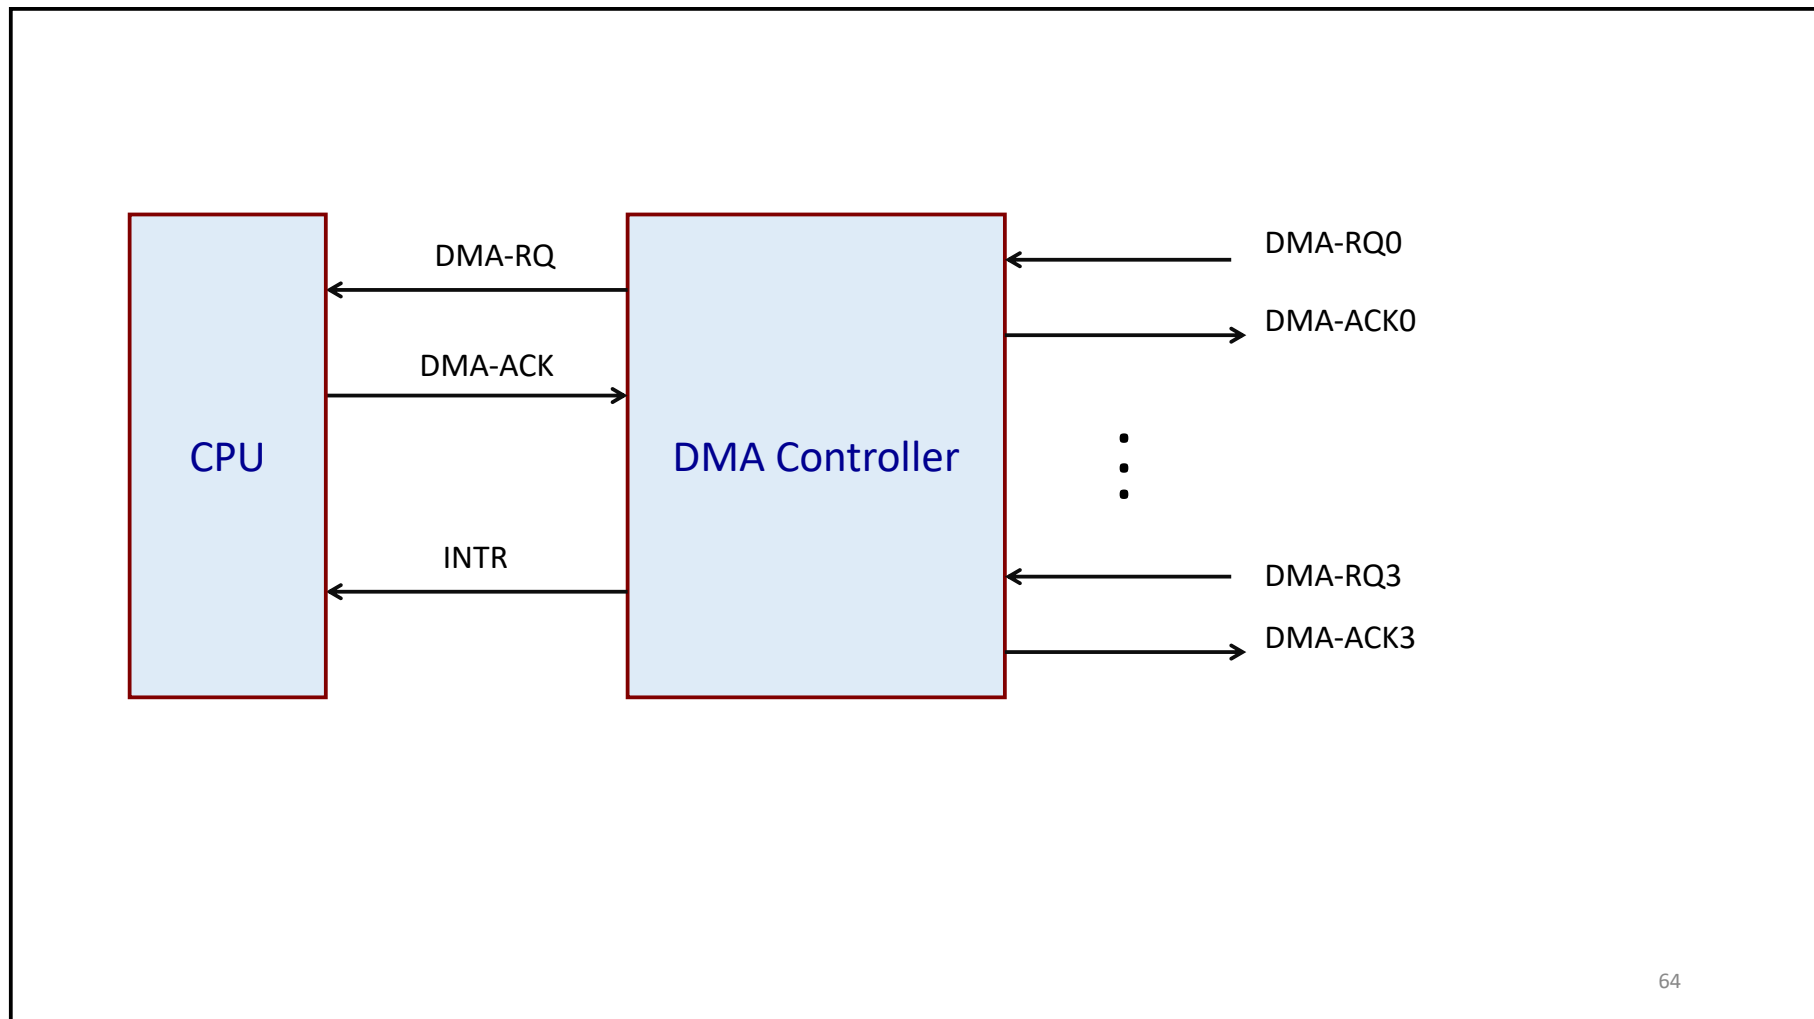


← Instruction Cycle →

*Why cannot we have
interrupt breakpoints at
the end of any machine
cycle?*

- For every DMA channel, the DMA controller will have three registers:
 - a) Memory address
 - b) Word count
 - c) Address of data on disk
- CPU initializes these registers before each DMA transfer operation.
- Before the data transfer, DMA controller requests the memory bus from the CPU.
- When the data transfer is complete, the DMA controller sends an interrupt signal to the CPU.

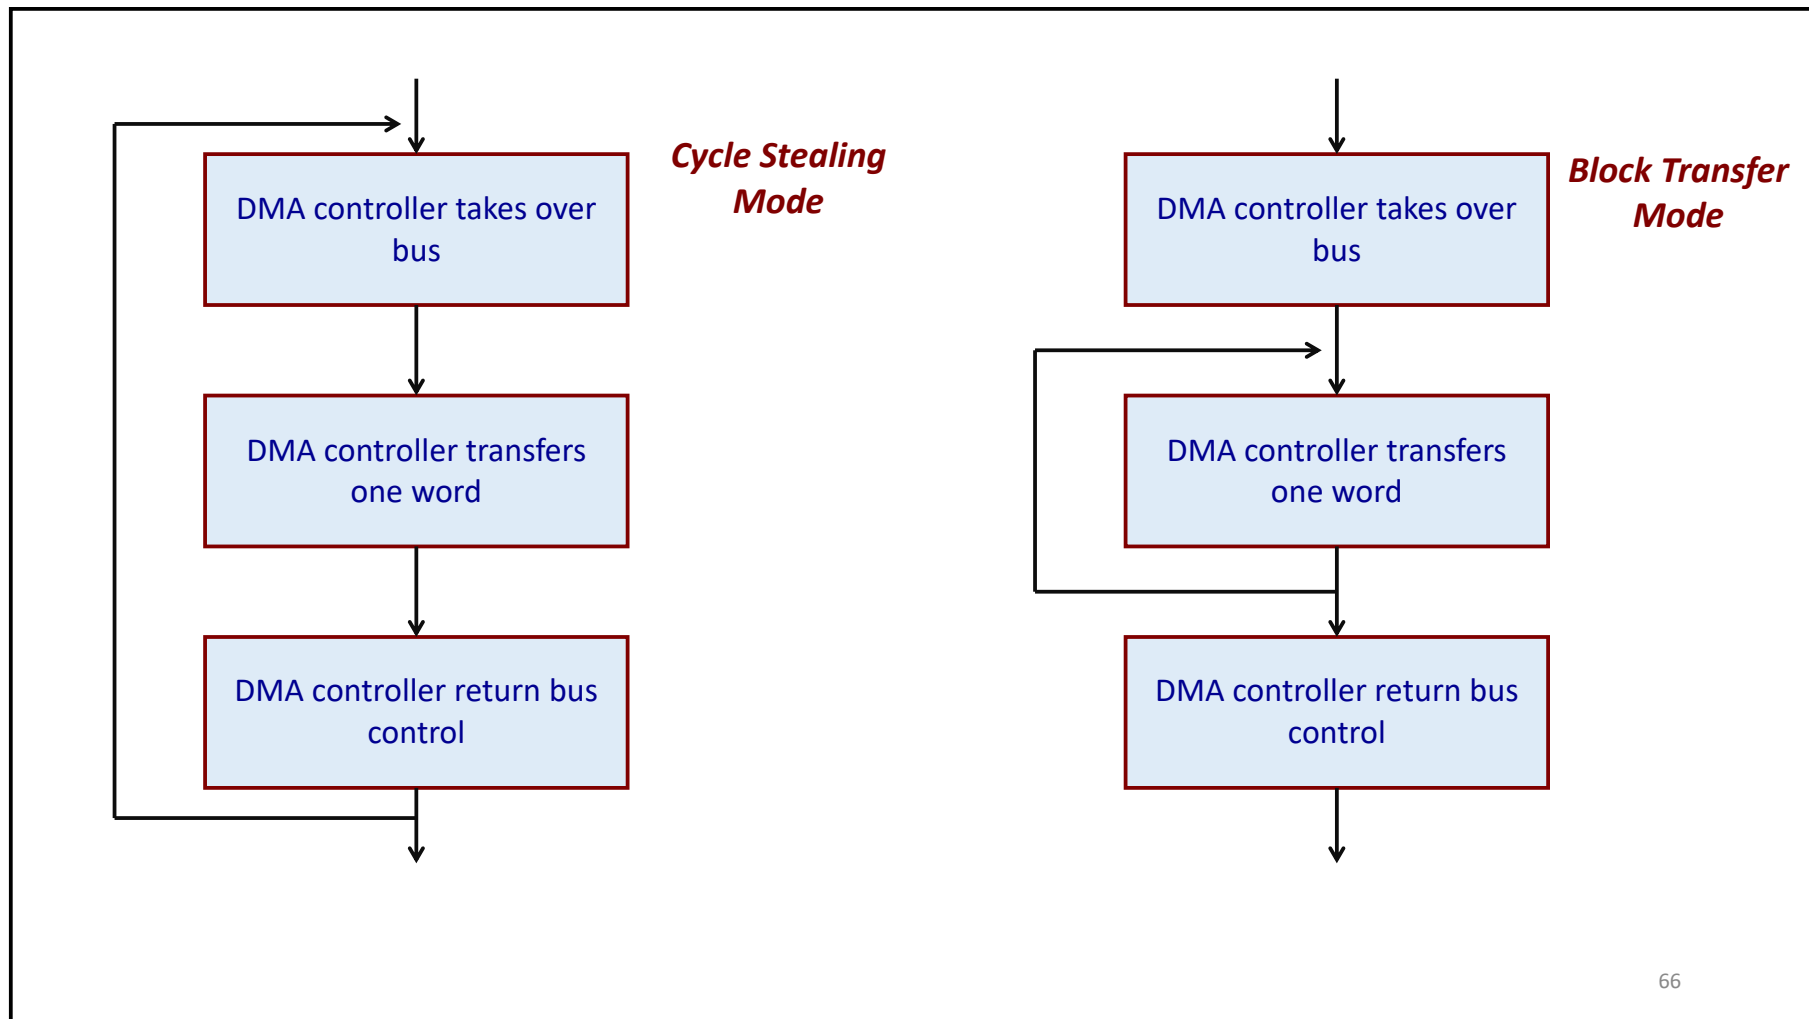
63



DMA Transfer Modes

- DMA transfer can take place in two modes:
 - a) **DMA cycle stealing**
 - The DMA controller requests for the for a few cycles 1 or 2.
 - Preferably when the CPU is not using memory.
 - DMA controller is said to steal cycles from the CPU without the CPU knowing it.
 - b) **DMA block transfer**
 - The DMA controller transfers the whole block of data without interruption.
 - Results in maximum possible data transfer rate.
 - CPU will lie idle during this period as it cannot fetch any instructions from memory.

65



Others Applications of DMA

- Other than data transfer to/from high-speed peripheral devices, DMA can be used in some other areas as well:
 - High-speed memory-to-memory block move.
 - Refreshing dynamic memory systems, by periodically generating dummy read requests to the columns.

67