

# CS60050

## Machine Learning

### Decision Trees: Overfitting and Pruning

Somak Aditya

Assistant Professor  
Department of CSE  
August 14, 2024

Sudeshna Sarkar

Professor  
CoE AI, Department of CSE



# Are decision trees algorithms optimal?

- Well, what do we mean by optimal?
- Considering all possible decision trees (i.e., trees splitting on one feature per node),
- will the ID3 algorithm (each split maximizes mutual information; stopping when mutual information is zero)...
- produce the **smallest** decision tree that has lowest **classification training error**?
- No, they aren't optimal
- Decision trees are **greedy algorithms**, i.e., they make the best local decision without considering longer term possibilities.
  - Better trees are possible, but it takes too long to search all combinations

# Decision Trees: Pros & Cons

- Pros

- Interpretable
- Efficient (computational cost and storage)
- Can be used for classification and regression tasks
- Compatible with categorical and real-valued features

## Cons

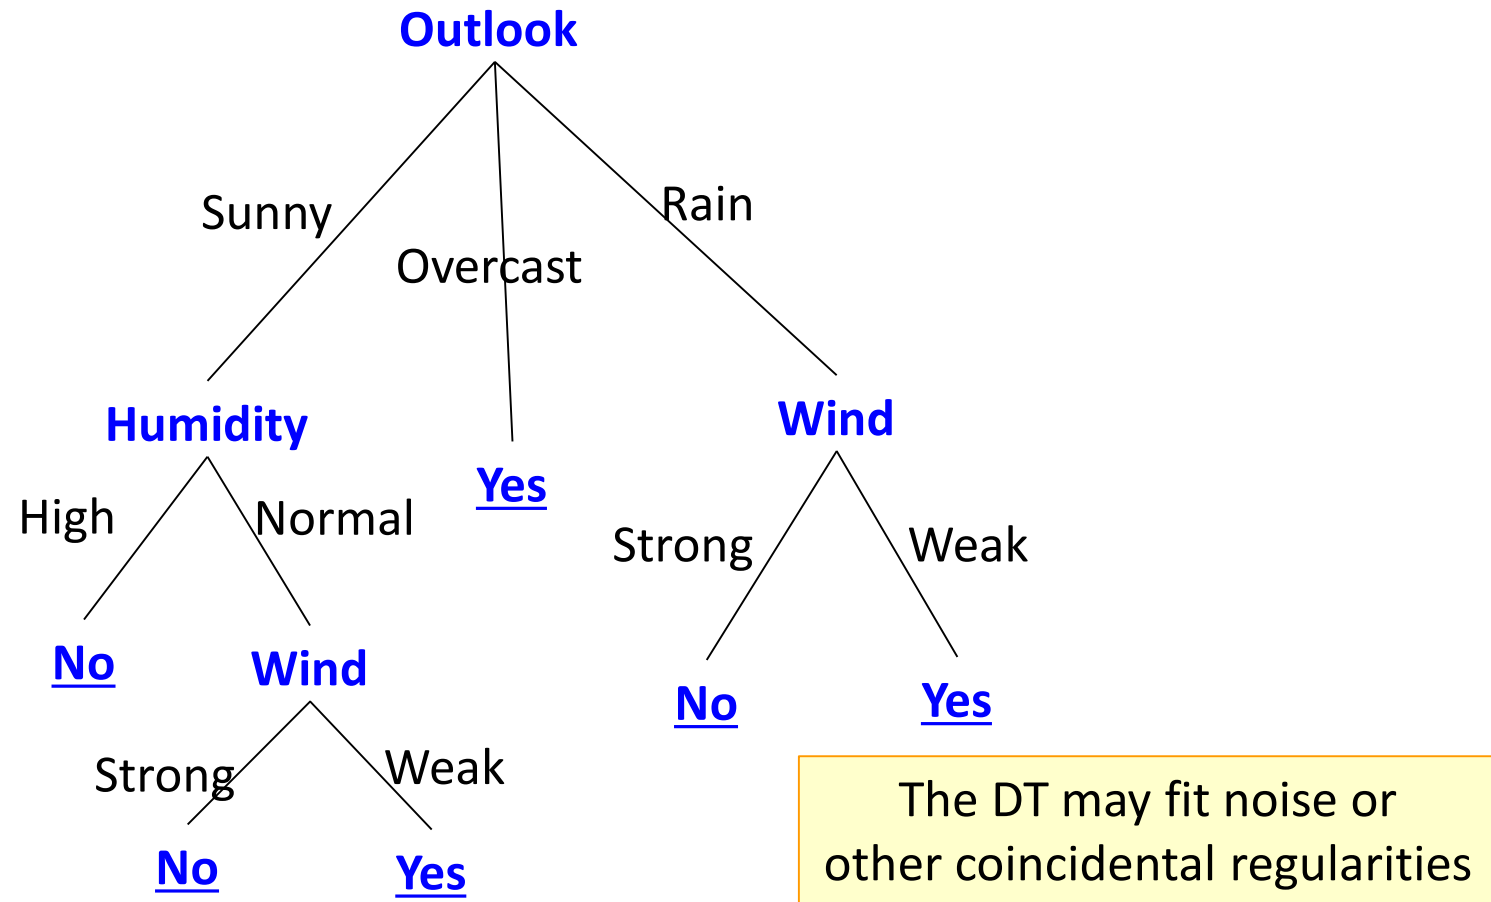
- Greedy: each split only considers the immediate impact
- Not guaranteed to find the smallest (fewest number of splits) tree that achieves a training error rate of 0.
- ***Liable to overfit!***

# Overfitting in Decision Trees

- Many kinds of “noise” can occur in the examples:
  - Two examples have same attribute/value pairs, **but different classifications**
  - Some values of attributes are incorrect **because of errors in the data acquisition process or the preprocessing phase**
  - The instance was **labeled incorrectly** (+ instead of -)
- Also, some attributes are irrelevant to the decision making process
  - e.g., color of a die is irrelevant to its outcome

# Overfitting - Example

Consider adding a **noisy** training example to the following tree:




What would be the effect of adding:

<Outlook = Sunny, Temp = Hot, Humidity = Normal, Wind = Strong, playTennis=NO>

# Training Data

Is (often)  
only a small  
set of the  
entire  
“instance  
space”



# Error Rate

Consider a hypothesis  $h$  over

- error over all training data:  $\text{error}(h, D_{\text{train}})$
- error rate over all test data:  $\text{error}(h, D_{\text{test}})$
- true error over all data:  $\text{error}_{\text{true}}(h, D)$

This is the quantity we care most about! But, in practice,  $\text{error}_{\text{true}}(h, D)$  is **unknown**.

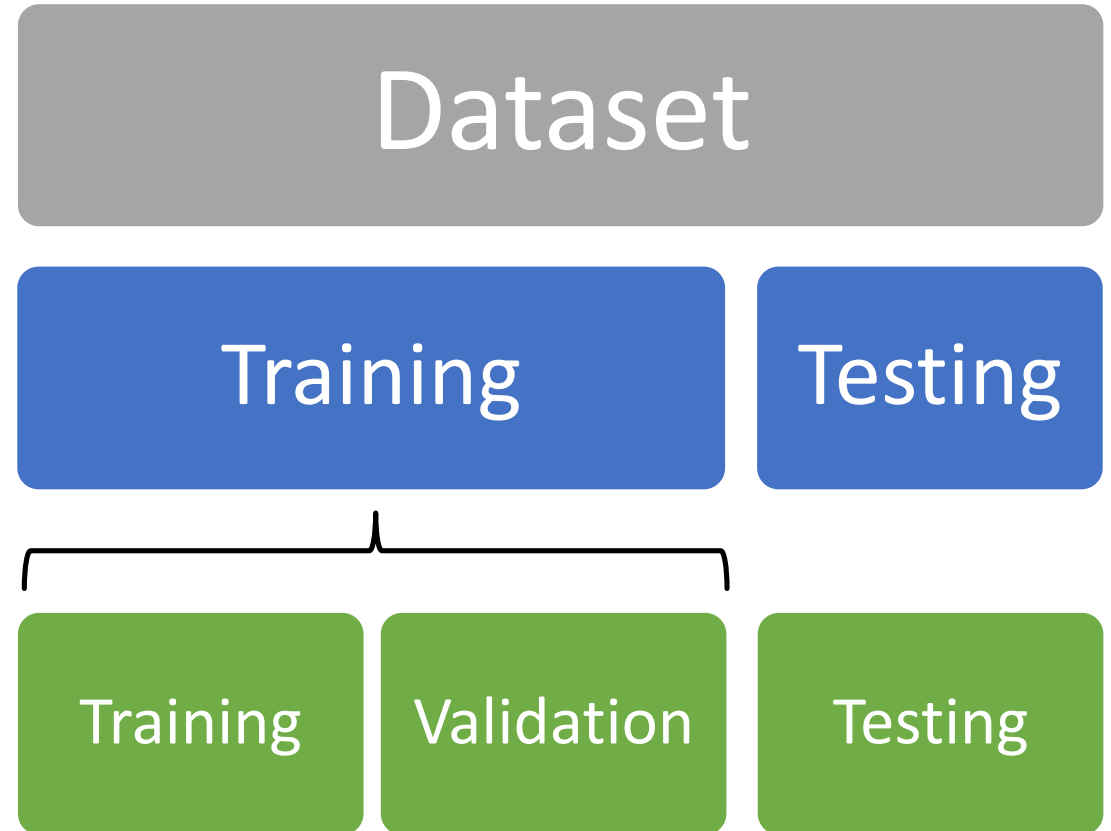
Learning a tree that classifies the training data perfectly may not lead to the tree with the *best generalization performance*.

- Noise in the training data
- Very little data

# Experimental Machine Learning

Split your data :


- Training data (e.g., 70-90%)
- Test data (e.g., 10-20%)
- Development data or Validation data (10-20%)
- You need to report performance on test data, but you are not allowed to look at it.
  - You are allowed to look at the development data (and use it to tune parameters)



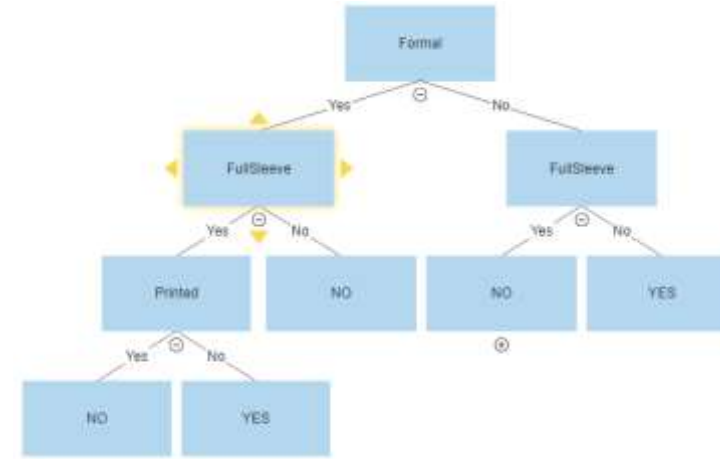
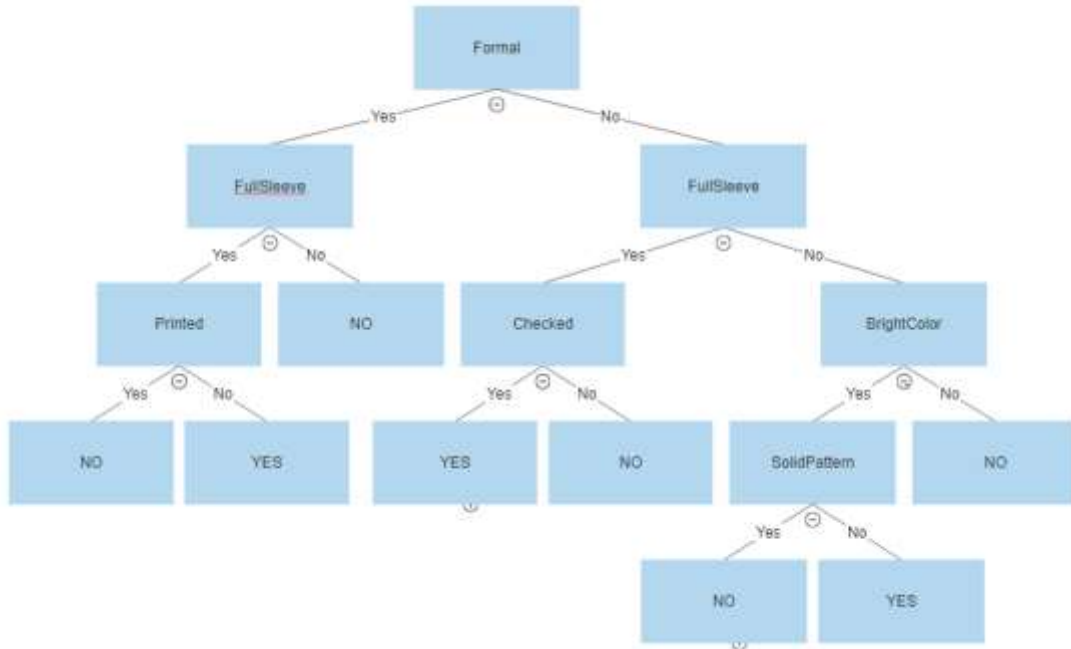


# Validation

- Divide your data randomly into training and *Validation* data.
- Build your best model based on the training data only.
- Apply your model to the Validation data.
- Does your model predict  $y'$  for the Validation data as well as it predicted  $y$  for the training data?



# Which Decision Tree?




# Overfitting

Overfitting :

- Fit the training data too well
- But fail to generalize to new examples

Why does Overfitting happen?

- Noise
- Irrelevant Features
- Insufficient Data
- Training data not representative



Overfitting results in decision trees that are more complex than necessary

# Overfitting



A hypothesis  $h$  is said to **overfit the training data** if there is another hypothesis  $h'$  such that  $h$  has smaller error than  $h'$  on the training data but  $h$  has larger error on the test data than  $h'$ .

In other words, hypothesis  $h$  overfits if there is  $h' \in \mathcal{H}$  such that

$$\text{error}_{\text{train}}(h) < \text{error}_{\text{train}}(h')$$


and

$$\text{error}_{\text{true}}(h) > \text{error}_{\text{true}}(h')$$




# Overfitting in Decision Trees

- Your model shows much greater loss on the test data than on the training data.
- **Example:** a decision tree with so many levels that the typical leaf is reached by only one member of the training set.



# Overfitting in Practice (ID3 – sklearn)

- ID3 performs heuristic search through space of decision trees
- It stops at smallest acceptable tree. Why?



**Occam's razor: prefer the simplest hypothesis that fits the data**

# Pruning a decision tree

- **Pruning The Tree:** remove unnecessary nodes to
  - make it more efficient and
  - solve overfitting problems.
- 1. **Prepruning:** Stop growing when data split not statistically significant
- 2. **Postpruning:** Grow full tree then remove nodes that seem not to have sufficient evidence.

Methods for evaluating subtrees to prune

- **Cross-validation: Reserve hold-out set to evaluate utility**
- Statistical testing: Test if the observed regularity can be dismissed as likely to occur by chance
- Minimum Description Length: Is the additional complexity of the hypothesis smaller than remembering the exceptions?

This is related to the notion of regularization – keep the hypothesis simple


# Avoid Overfitting

- How can we avoid overfitting a decision tree?
  - **Prepruning:** Stop growing when data split not statistically significant
  - **Postpruning:** Grow full tree then remove nodes



# Pre-Pruning (Early Stopping)

- Early Stopping: Stop the learning algorithm before tree becomes too complex



## Stopping conditions:

- Do not split a node which contains too few instances
- Stop if expanding the current node does **not improve impurity measures significantly** (e.g., Gini or information gain)
- Limit tree depth

# Reduced-error Pruning

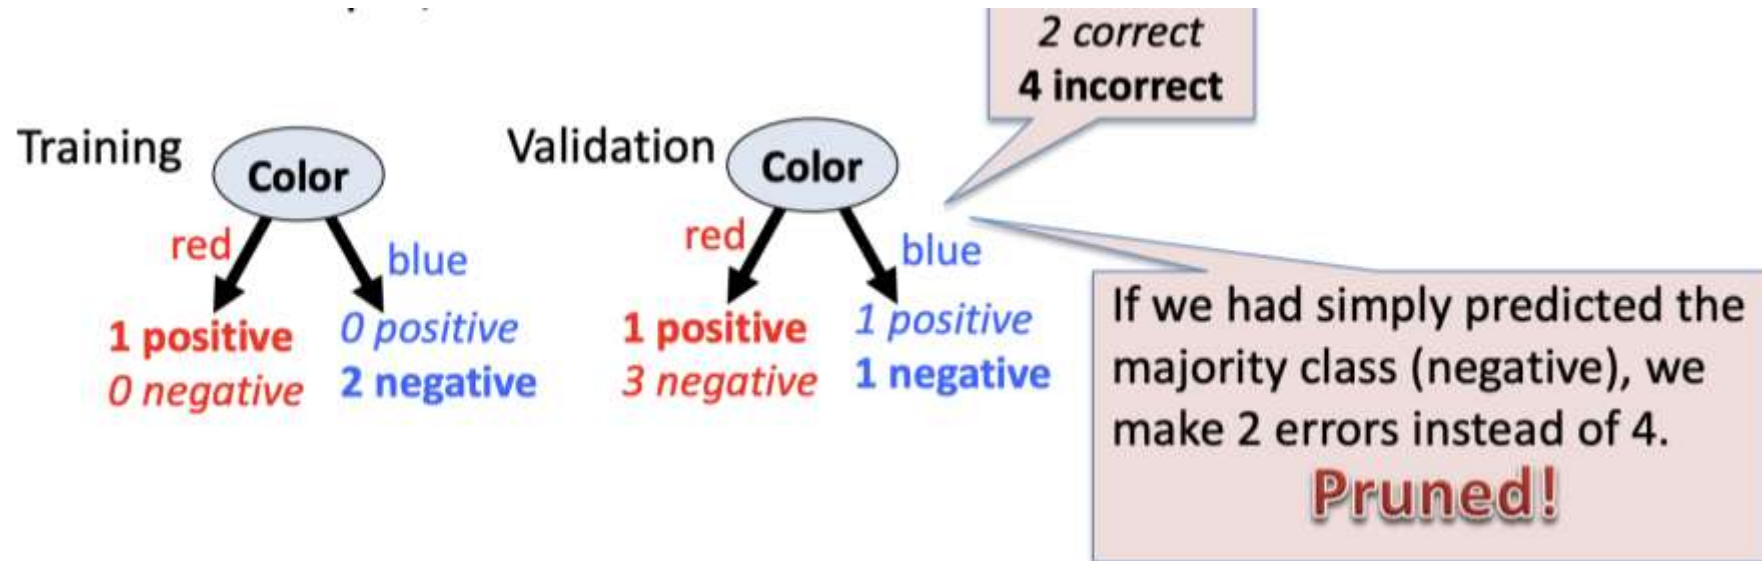
Partition data into train set and validation set

- Build a tree using the train set.
- Until accuracy on validation set decreases, do:
  - For each non-leaf node in the tree
    - ✓ Temporarily prune the tree below; replace it by majority vote
    - ✓ Test the accuracy of the hypothesis on the validation set
    - ✓ Permanently prune the node with the greatest increase in accuracy on the validation test.


# Pruning Decision Trees

Pruning the decision tree is done by replacing a whole subtree by a leaf node. The replacement takes place if a decision rule establishes that

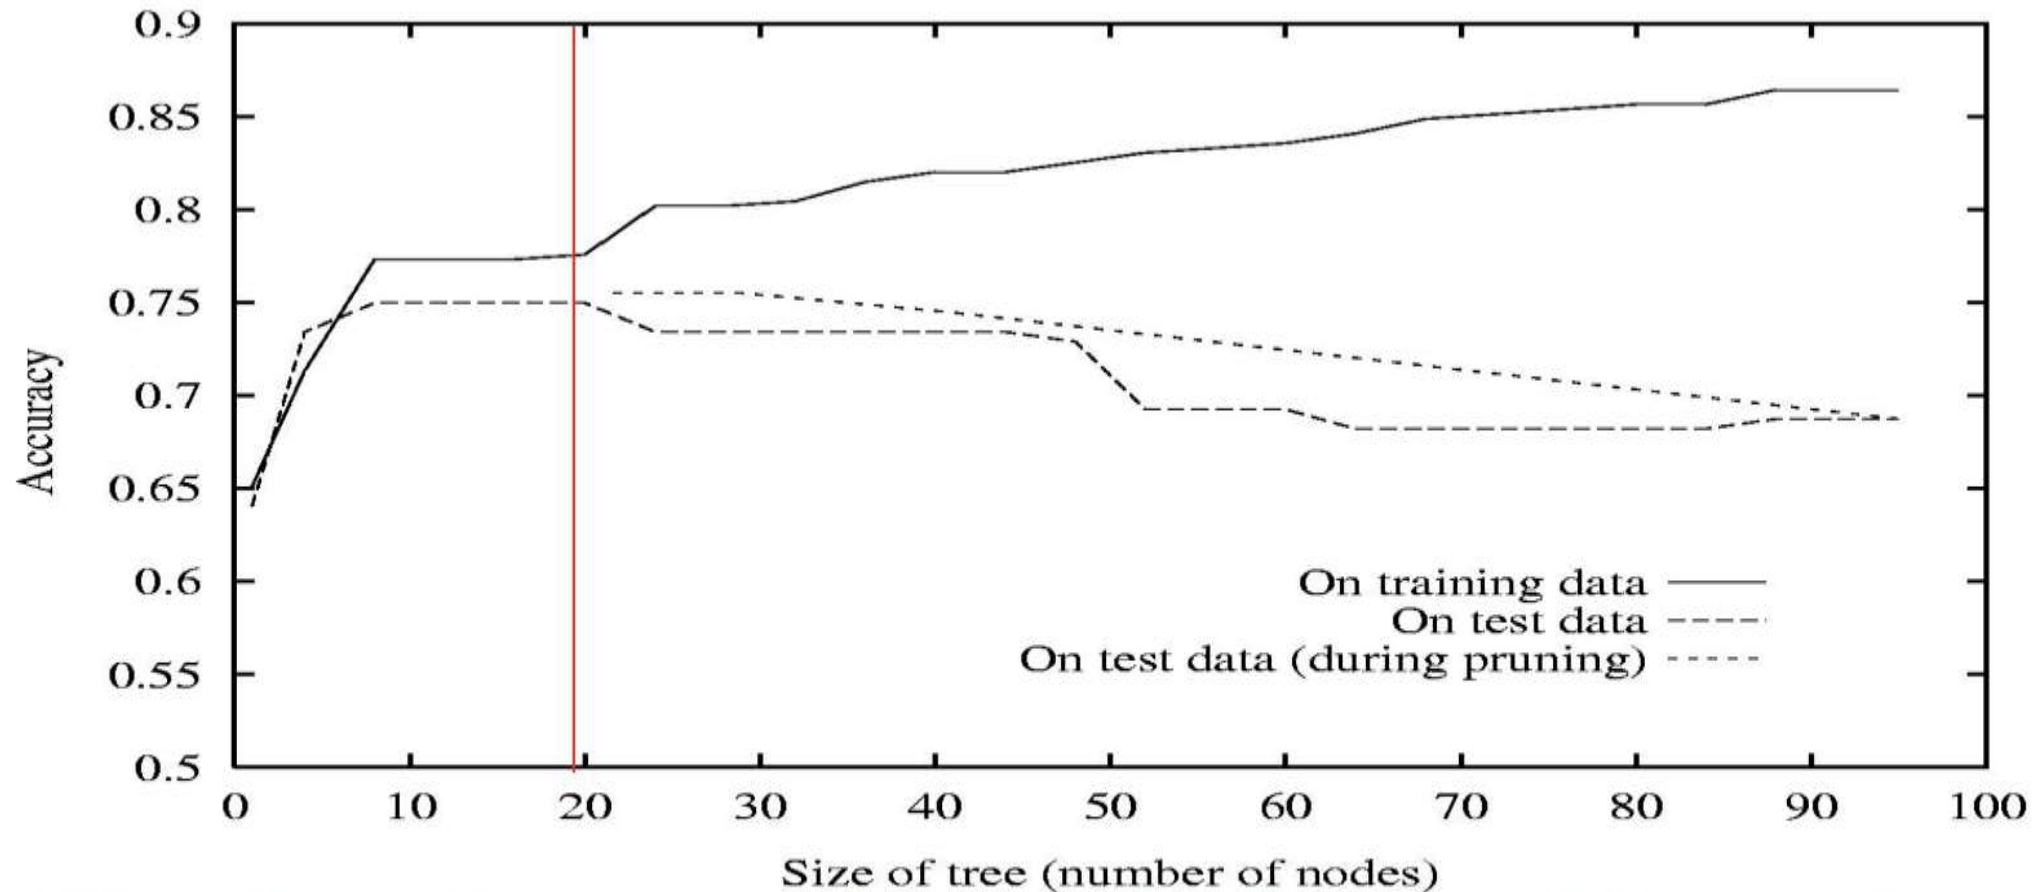
- the Expected Error Rate in the subtree  $>$  Expected error rate in the single leaf
- For example



# Effect of Reduced Error Pruning



# Effect of Reduced-Error Pruning



**The tree is pruned back to the red line where it gives more accurate results on the test data**

# Bias & Variance


# Overfitting vs Underfitting

## Underfitting

- Not able to capture the concept
  - Features don't capture concept
  - Model is not powerful.

## Overfitting


- Fitting the data too well



# Function Approximation: The Big Picture

Instance Space  $\mathcal{X} = \{0, 1\}^d$   
 $\mathbf{x} = \langle x_1, x_2, \dots, x_d \rangle \in \mathcal{X}$

Hypothesis Space  
 $H = \{h \mid h : \mathcal{X} \mapsto \{0, 1\}\}$

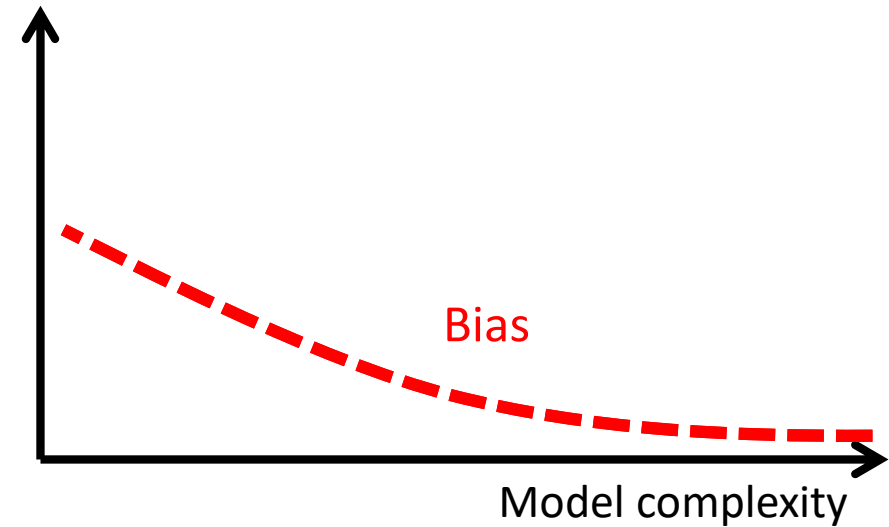


- How many labeled instances are needed to determine which of the  $2^{2^{20}}$  hypotheses are correct?
  - All  $2^{20}$  instances must be labeled!
- Generalizing beyond the training data (inductive inference) is impossible unless we add more assumptions (e.g., priors over  $H$ )



# Bias of a Learner ( $\sim$ mean error)

- How likely is the learner to identify the **target** hypothesis?
- Bias is **low** when the model is expressive (low empirical error)
- Bias is **high** when the model is too simple
  - The larger the hypothesis space is, the easiest it is to be close to the true hypothesis.
  - For each data set  $D$ ,
    - You learn a different hypothesis  $h(D)$ , that has a different true error  $error_{true}(h)$ ;
    - difference of the mean of this random variable from the true error.




if we train models  $f_D(X)$  on many training sets  $D$ , bias is the expected difference between their predictions and the true  $y$ 's.

$$Bias = E[f_D(X) - y]$$

# Variance of a Learner

How susceptible is the learner to different subsets of the training data? (i.e. to **different**  $D \sim P(\mathbf{X}, Y)$  )




Variance increases with model complexity

- **The larger the hypothesis space**, the more flexible the selection of the chosen hypothesis is as a function of the data.
- For each data set  $D$ ,
  - you will learn a different hypothesis  $h(D)$ , that will have a different error  $error_{true}(h)$ ;
  - Lets see the variance of this random variable.

if we train models  $f_D(X)$  on many training sets  $D$ , the variance of the estimates:

$$Variance = E \left[ \left( f_D(X) - \bar{f}(X) \right)^2 \right] \quad (\sim \text{std.dev among predictions})$$

# Impact of bias and variance



Expected error  $\approx$  bias + variance (why???)

# Bias-Variance Decomposition of Squared Error

- Assume that  $y = f(x) + \epsilon$ 
  - Noise  $\epsilon$  is sampled from a normal distribution with 0 mean and variance  $\sigma^2$ :  $\epsilon \sim N(0, \sigma^2)$
  - Noise lower-bounds the performance (error) we can achieve.

- Recall the objective function:

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( y^{(i)} - h_{\theta}(x^{(i)}) \right)^2$$

- We view this as an approximation of the expected value of the squared error:  $E(y - h_{\theta}(x))^2$

# Bias-Variance Decomposition of Squared Error

$$\begin{aligned} E(y - h_{\theta}(x))^2 &= E[(y - f(x) + f(x) - h_{\theta}(x))^2] \\ &= E[(y - f(x))^2] + E[(f(x) - h_{\theta}(x))^2] + 2E[(f(x) - h_{\theta}(x))(y - f(x))] \\ &= E[(y - f(x))^2] + E[(f(x) - h_{\theta}(x))^2] + 2(\cancel{E[f(x)h_{\theta}(x)]} + \cancel{E[yf(x)]} \\ &\quad - \cancel{E[yh_{\theta}(x)]} - \cancel{E[f(x)^2]}) \end{aligned}$$

Therefore

$$\begin{aligned} E(y - h_{\theta}(x))^2 &= E[(y - f(x))^2] + E[(f(x) - h_{\theta}(x))^2] \\ &= E[\epsilon^2] + E[(f(x) - h_{\theta}(x))^2] \end{aligned}$$

**Aside:**


**Definition of Variance**

$$\text{var}(z) = E[(z - E[z])^2]$$

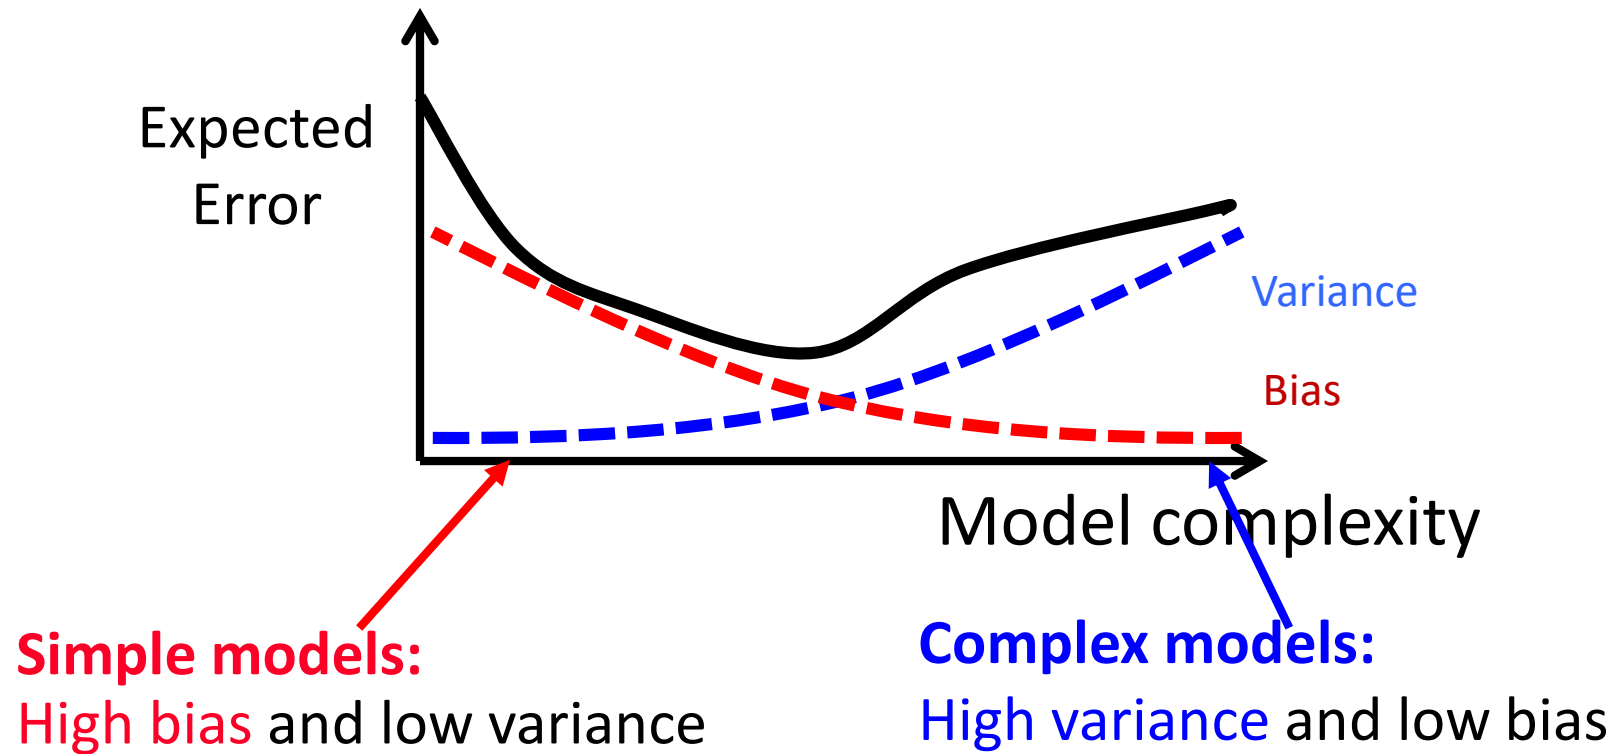
This is  $\text{var}(\epsilon)$  since mean is 0.

# Bias-Variance Decomposition of Squared Error

$$\begin{aligned} \mathbb{E}[(y - h_{\theta}(\mathbf{x}))^2] &= \text{var}(\epsilon) + \mathbb{E}[(f(\mathbf{x}) - h_{\theta}(\mathbf{x}))^2] \\ &= \text{var}(\epsilon) + \mathbb{E}[(f(\mathbf{x}) - \mathbb{E}[h_{\theta}(\mathbf{x})] + \mathbb{E}[h_{\theta}(\mathbf{x})] - h_{\theta}(\mathbf{x}))^2] \\ &= \text{var}(\epsilon) + \mathbb{E}[(f(\mathbf{x}) - \mathbb{E}[h_{\theta}(\mathbf{x})])^2] + \mathbb{E}[(\mathbb{E}[h_{\theta}(\mathbf{x})] - h_{\theta}(\mathbf{x}))^2] \\ &\quad + 2\mathbb{E}[(\mathbb{E}[h_{\theta}(\mathbf{x})] - h_{\theta}(\mathbf{x}))(f(\mathbf{x}) - \mathbb{E}[h_{\theta}(\mathbf{x})])] \\ &= \text{var}(\epsilon) + \mathbb{E}[(f(\mathbf{x}) - \mathbb{E}[h_{\theta}(\mathbf{x})])^2] + \mathbb{E}[(\mathbb{E}[h_{\theta}(\mathbf{x})] - h_{\theta}(\mathbf{x}))^2] \\ &\quad + 2(\cancel{\mathbb{E}[f(\mathbf{x})\mathbb{E}[h_{\theta}(\mathbf{x})]]} - \cancel{\mathbb{E}[\mathbb{E}[h_{\theta}(\mathbf{x})]^2]} - \cancel{\mathbb{E}[f(\mathbf{x})h_{\theta}(\mathbf{x})]} + \cancel{\mathbb{E}[h_{\theta}(\mathbf{x})\mathbb{E}[h_{\theta}(\mathbf{x})]]}) \end{aligned}$$

  
cancels                  cancels

# Model complexity



# BIAS

- Error caused because the model can not represent the concept
- Bias is the expected difference between the model prediction and the true  $y$ 's.
- **Higher Bias:**
  - Decision tree of lower depth
  - Linear functions
  - Important features missing

if we train models  $f_D(X)$  on many training sets  $D$ , bias is the expected difference between their predictions and the true  $y$ 's.

$$Bias = E[f_D(X) - y]$$

# VARIANCE

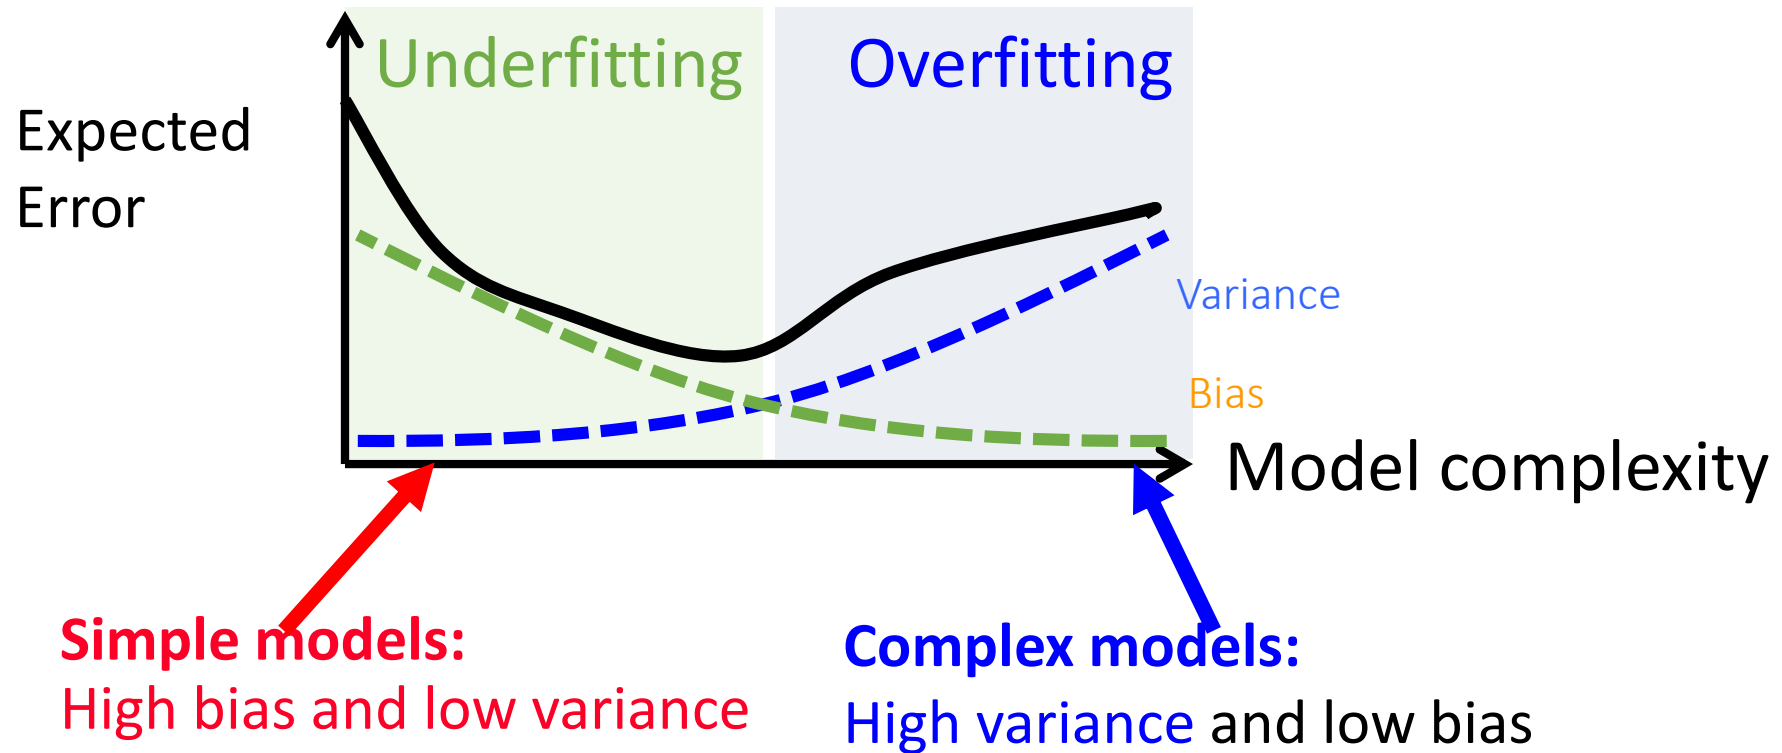
- Error caused because the learned model reacts to small changes (noise) in the training data
- High variance can cause an algorithm to model the random noise in the training data, rather than the intended outputs
- **Higher Variance**
  - Decision tree with large no of nodes
  - High degree polynomials
  - Many features

if we train models  $f_D(X)$  on many training sets  $D$ , variance is the variance of the estimates:

$$Variance = E\left[\left(f_D(X) - \bar{f}(X)\right)^2\right]$$



# Underfitting and Overfitting



This can be made more accurate for some loss functions.


We will discuss a more precise and general theory that trades **expressivity of models** with **empirical error**

# Bias and Variance Tradeoff

There is usually a bias-variance tradeoff caused by model complexity.

**Complex models** often have lower bias, but higher variance.

**Simple models** often have higher bias, but lower variance.




# Trade-Offs

$$\text{Error} \approx \text{Function}(\text{Complexity}, \text{TrainingDataSize})$$

- There is a trade-off between these factors:
    - Complexity of Model  $c(H)$
    - Training set size,  $m$ ,
    - Generalization error,  $E$  on new data
1. As  $m$  *increases*,  $E$  *decreases*
  2. As  $c(H)$  *increases*,
    1. first  $E$  *decreases* and then  $E$  *increases*
    2. the training error *decreases* for some time and then stays constant (frequently at 0)

As  $m$  increases,  $E$  decreases



# Model complexity

2. As  $c(H)$  increases, first  $E$  decreases and then  $E$  increases
3. As  $c(H)$  *increases*, the training error *decreases* for some time and then stays constant (frequently at 0)

